

COMP5338 – Advanced Data Models

Week 3: MongoDB – Aggregation Framework

Dr. Ying Zhou
School of Computer Science



Outline

■ Null type

■ Aggregation

- ▶ Single collection aggregation
- ▶ Aggregation pipeline with multiple collection

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

WARNING

This material has been reproduced and communicated to you by or on behalf of the **University of Sydney** pursuant to Part VB of the Copyright Act 1968 (the Act).

The material in this communication may be subject to copyright under the Act. Any further reproduction or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice

Null, empty string and related operators

- Null (or null) is a special data type
 - ▶ Similar to **None**, **Null** or **Nil** in any programming language
 - ▶ It has a singleton value expressed as **null**
 - ▶ Indicating no value is here
- The interpretation of null is different depending on where it appears
- It might represents
 - ▶ The field exists, but has no value
 - ▶ The field does not exists
 - ▶ Or both
- This is different to giving a field an empty string "" as value

<https://docs.mongodb.com/manual/tutorial/query-for-null-fields/index.html>

Null query example

■ Collection revisions document sample

```
{ "_id" : ObjectId("5799843ee2cbe65d76ed919b"),  
  "title" : "Hillary_Clinton",  
  "timestamp" : "2016-07-23T02",  
  "revid" : 731113635,  
  "user" : "BD2412",  
  "parentid" : 731113573,  
  "size" : 251742,  
  "minor" : "" }
```

```
1 {  
2   "_id" : ObjectId("5d42869d0c84336545f9b2d3"),  
3   "title" : "Hillary_Clinton",  
4   "timestamp" : ISODate("2016-07-01T19:33:39.000Z"),  
5   "parsedcomment" : "The religious affiliation of Hillary",  
6   "revid" : 727871391,  
7   "user" : "Theologicalmess",  
8   "parentid" : 727749878,  
9   "size" : 246162  
10 }
```

- We need a field to indicate if a revision is *minor* or not. The original schema uses a field with empty string value to indicate a minor revision; a document without this field would be a non-minor revision or major revision.

<https://docs.mongodb.com/manual/tutorial/query-for-null-fields/>

Querying for null or field existence

■ Queries

- ▶ `db.revisions.find({minor:{$exists:true}})`
 - Find all documents that has a field called **minor**
- ▶ `db.revisions.find({minor:""})`
 - Find all documents whose **minor** field has a value of "", empty string
- ▶ `db.revisions.find({minor: {$ne: null}})`
 - Find all documents whose **minor** field's value is not **null**
- ▶ `db.revisions.find({minor:null})`
 - Find all documents that do not have a **minor** field or the value of **minor** field is null
- ▶ `db.revisions.find({minor:{$exists:false}})`
 - Find all documents that does not have a **minor** field

It is possible to set the value to null

```
db.revisions.insertOne({title:"nulltest",  
  "timestamp" : "2018-08-14T02:02:06Z",  
  "revid" : NumberLong(7201808141159),  
  "user" : "BD2412",  
  "parentid" : 731113573,  
  "size" : NumberInt(251900),  
  "minor":null})
```

```
db.revisions.insertOne({title:"nulltest",  
  "timestamp" : "2018-08-14T02:02:06Z",  
  "revid" : NumberLong(201808141157),  
  "user" : "BD2412",  
  "parentid" : NumberLong(731113573),  
  "size" : NumberInt(251800)})
```

`db.revisions.find({minor:null})` would return both documents

`db.revisions.find({minor:{$exists:false}})` can differentiate the two

Outline

- Null type

- Aggregation

- ▶ Single collection aggregation pipeline
- ▶ Aggregation pipeline with multiple collection

Aggregation

- Simple and relatively standard data analytics can be achieved through **aggregation**
 - ▶ Grouping, summing up value, counting, sorting, etc
 - ▶ Running on the DB engine instead of application layer
- Several options
 - ▶ Aggregation Pipeline
 - ▶ MapReduce
 - Through JavaScript Functions
 - Is able to do customized aggregations

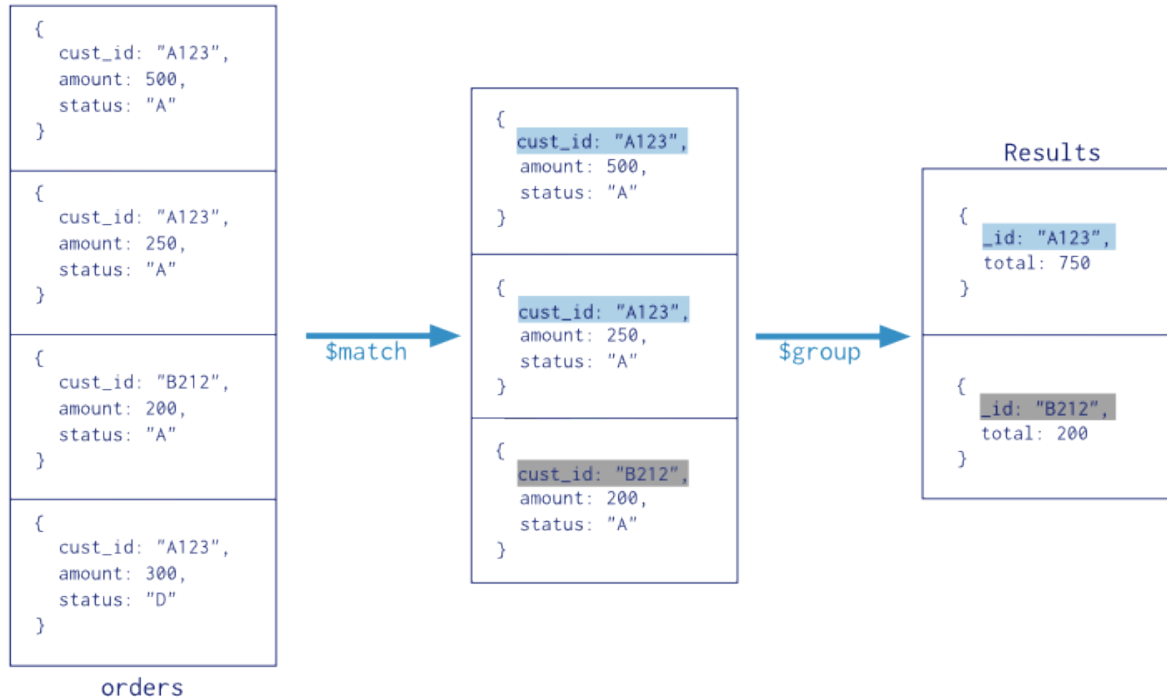
Aggregation Pipeline

- Aggregation pipeline consists of multiple stages
 - ▶ Each stage transforms the incoming documents as expressed in the stage object
 - ▶ The stage object is enclosed in a pair of curly braces
 - ▶ The pipeline is an array of many stage objects.

```
db.collection.aggregate( [  
    { stage name: {expression,..., expression} },  
    { stage name: {expression,..., expression} },  
    ...  
] )
```

Aggregation Example

Collection
↓
`db.orders.aggregate([`
 `$match stage → { $match: { status: "A" } },`
 `$group stage → { $group: { _id: "$cust_id", total: { $sum: "$amount" } } }`
 `]` `)`



```
select  cust_id as _id, SUM(amount) as total
from    orders
where   status = "A"
group by cust_id
```

Typical aggregation stages

- **\$match**
- **\$group**
- **\$project**
- **\$unwind**
- **\$sort**
- **\$skip**
- **\$limit**
- **\$count**
- **\$sample**
- **\$out**
- **\$lookup**

\$match stage

- **\$match**: filters the incoming documents based on given conditions

- **Format:**

`{ $match: {<query> } }`

- The query document is the same as those in the **find** query

- **Example:**

```
db.revisions.aggregate([ { $match: { size : { $lt: 250000 } } } ] )
```

Has the same effect as

```
db.revisions.find({ size : { $lt: 250000 } })
```

\$group stage

- **\$group**: groups incoming documents by some specified expression and outputs to the next stage a document for each distinct group

- **Format:**

```
{ $group:
  { _id: <expression>, // Group By Expression
    <field1>: { accumulator: <expression> },
    ... }
}
```

- ▶ The `_id` field of the output document has the value of the group key for each group
- ▶ The other fields usually represent the statistics you want to produce for each group
- ▶ One statistics per field
 - Total amount, average price, group size

\$group stage (cont'd)

- **<Expression>** in **{_id:<expression>}**,
 - ▶ **null** value, to specify the whole collection as a group
 - ▶ field path to specify one or many fields as grouping key
 - Field name prefixed with \$ sign in a pair of quotes
 - “\$title”, or “\$address.street”
- **{accumulator: <expression>}**
 - ▶ There are predefined accumulators: \$sum, \$avg, \$first, \$last, etc
 - ▶ User defined accumulators can be used as well
 - ▶ Field path will be used in the <expression> if the accumulator returns value based on field values in the incoming document

\$group stage example

- Find the earliest revision time in the whole collection

```
db.revisions.find({}, {timestamp:1, _id:0})
```

```
.sort({timestamp:1})
```

 Sort in ascending order of timestamp

```
.limit(1)
```

```
db.revisions.aggregate([  
  {$group: {_id:null, earliest: {$min: "$timestamp"}}}  
])
```

Accumulator Field path as expression

Returns a single document

- Find the earliest revision time of each page in the collection

```
db.revisions.aggregate([  
  {$group: {_id:"$title", earliest: {$min: "$timestamp"}}}  
])
```

field path

Returns a document for each distinct title

\$group stage example (cont'd)

- Find the number of revisions made on each page by each individual user
 - ▶ This would require grouping based on two fields: title and user
 - ▶ We need to specify these two as the `_id` field of the output document

Composite type as `_id`

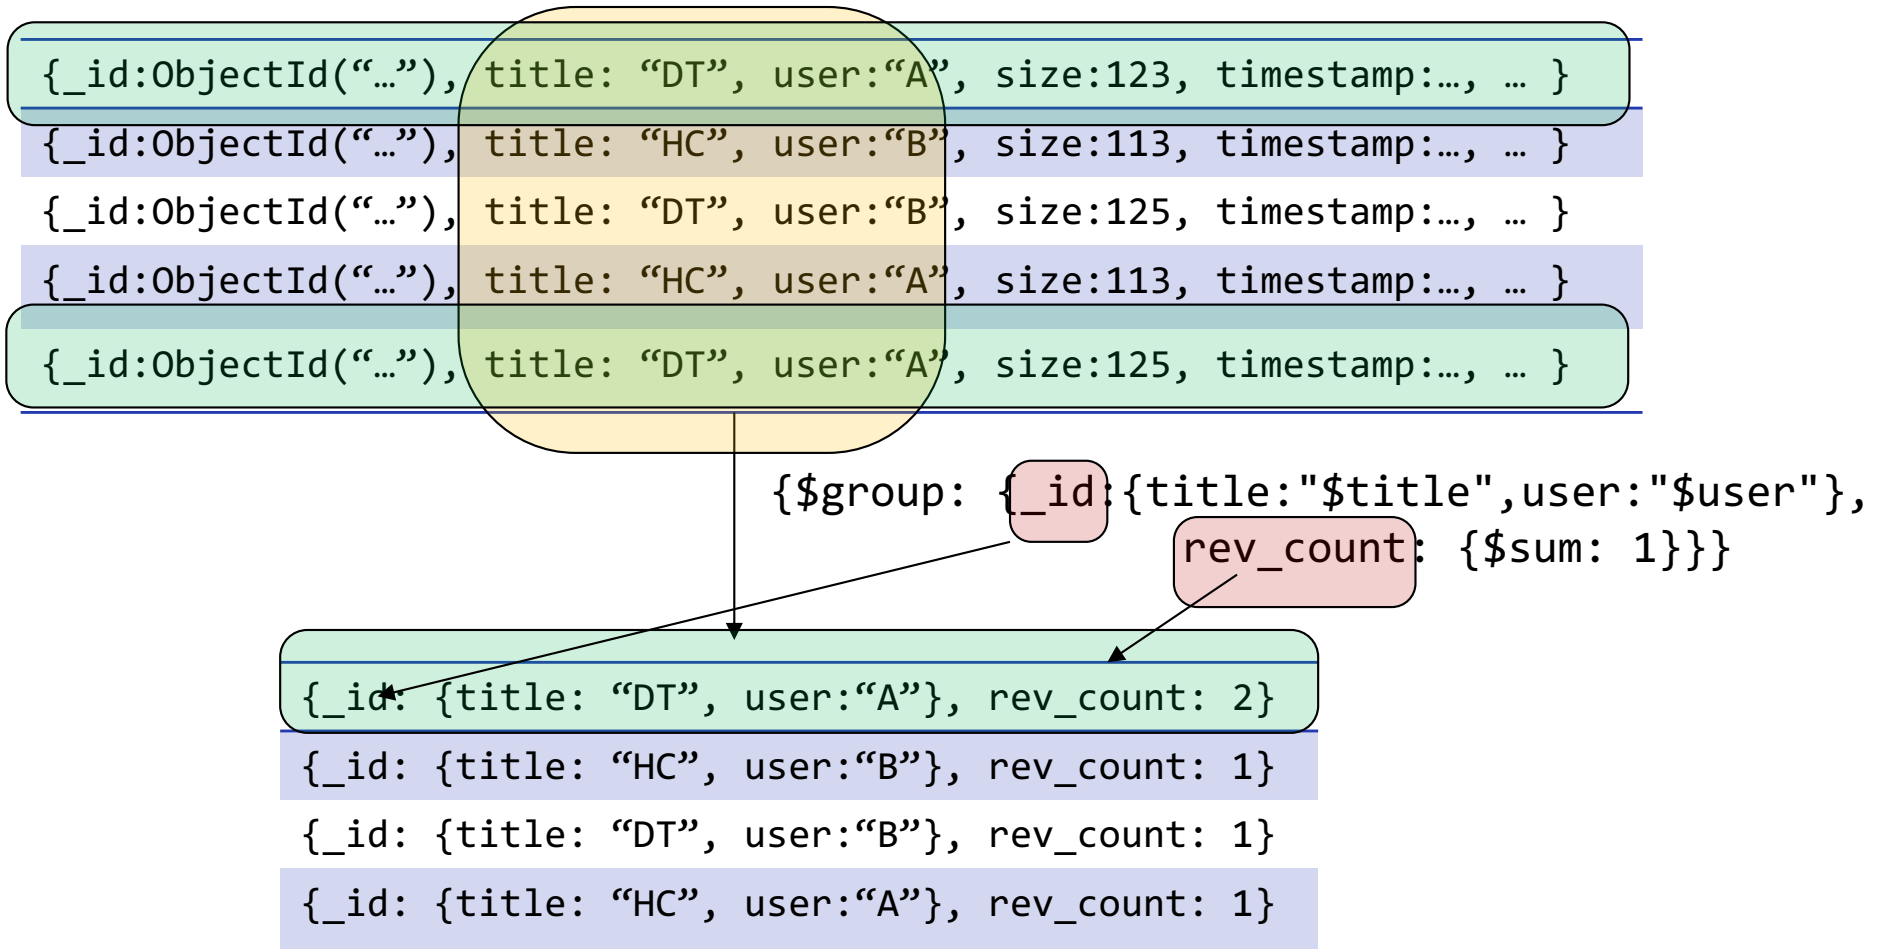
```
db.revisions.aggregate([  
  {$group: {_id: {$title: "$title", user: "$user"},  
    rev_count: {$sum: 1}}}  
])
```

accumulator

number literal

Same effect as count

\$group by more than one field



\$group examples (cont'd)

- Accumulators do not *just* return a single value, we can use accumulators to create an array to hold data from incoming documents
- Example of two commands:

```
db.revisions.aggregate([
  {$group: {_id: "$title",
            revs: {$push: {user: "$user", timestamp: "$timestamp"}}}}
])
```

```
db.revisions.aggregate([
  {$group: {_id: "$title", rev_users: {$addToSet: "$user"}}}}
])
```

They have the same group key: \$title

They have another field in addition to the group key

The other field is created with different accumulators

\$push accumulator

```
{_id:ObjectId("..."), title: "DT", user:"A", size:123, timestamp:..., ... }
{_id:ObjectId("..."), title: "HC", user:"B", size:113, timestamp:..., ... }
{_id:ObjectId("..."), title: "DT", user:"B", size:125, timestamp:..., ... }
{_id:ObjectId("..."), title: "HC", user:"A", size:113, timestamp:..., ... }
{_id:ObjectId("..."), title: "DT", user:"A", size:125, timestamp:..., ... }
```

↓

```
db.revisions.aggregate([
  {$group:
    {_id:"$title",
     revs:{$push:{user:"$user",timestamp:"$timestamp"}}}
  ]})
```

```
{ _id: "DT",
  revs:[
    {user:"A",timestamp:...},
    {user:"B",timestamp:...},
    {user:"A",timestamp:...}
  ]}
```

```
{ _id:"HC",
  revs:[
    {user:"A", timestamp:...},
    {user:"B", timestamp:...}
  ]}
```

\$addToSet accumulator

```
{_id:ObjectId("..."), title: "DT", user:"A", size:123, timestamp:..., ... }  
{_id:ObjectId("..."), title: "HC", user:"B", size:113, timestamp:..., ... }  
{_id:ObjectId("..."), title: "DT", user:"B", size:125, timestamp:..., ... }  
{_id:ObjectId("..."), title: "HC", user:"A", size:113, timestamp:..., ... }  
{_id:ObjectId("..."), title: "DT", user:"A", size:125, timestamp:..., ... }
```

```
db.revisions.aggregate([  
  {$group: {_id:"$title",  
            rev_users:{$addToSet:"$user"}}}  
])
```

```
{ _id: "DT",  
  rev_users:["A", "B"]  
}
```

```
{ _id:"HC",  
  rev_users:["A", "B"]  
}
```

\$project stage

■ \$project

- ▶ **Restructure** the document by including/excluding field, adding new fields, resetting the value of existing field
- ▶ More powerful than the *project* argument in **find** query
- ▶ Format

{*\$project*: {<specification(s)>}}

- ▶ The specification can be an existing field name followed by a single value indicating the inclusion (1) or exclusion (0) of fields
- ▶ Or it can be a field name (existing or new) followed by an expression to compute the value of the field

<field>: <expression>

- ▶ In the expression, existing field from incoming document can be accessed using field path: “*\$fieldname*”

\$project examples

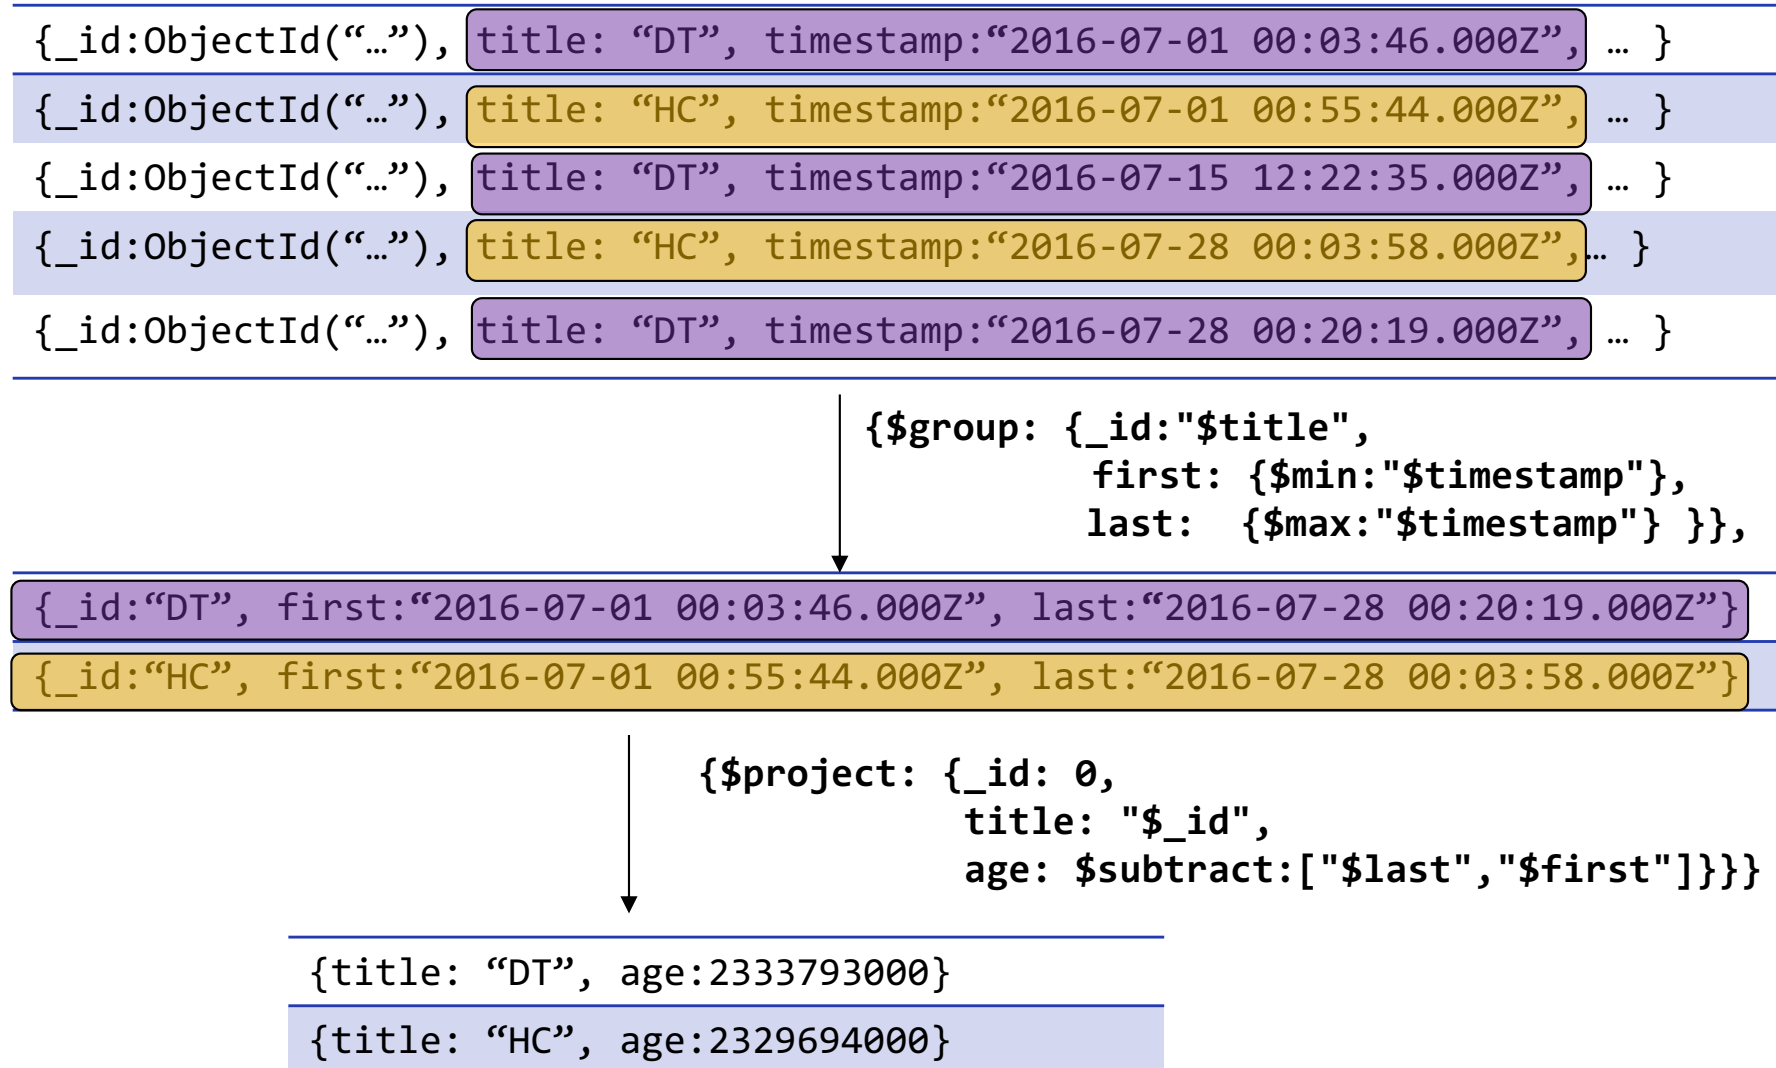
- Find the **age** of each title in the collection, where the **age** is defined as the duration between the last and the first revision of that title, assuming the timestamp is of ISODate type

```
db.revisions.aggregate([
  {$group: {_id:"$title",
            first: {$min:"$timestamp"},
            last:  {$max:"$timestamp"} }},
  {$project: {_id: 0,
              title: "$_id",
              age: {$subtract:["$last","$first"]}}}
])
```

Arithmetic expression operator, part of a large group of **Aggregation pipeline operator**

<https://docs.mongodb.com/manual/reference/operator/aggregation/#arithmetic-expression-operators>

\$group then \$project



We can combine multiple operators

```
db.revisions.aggregate([
  {$group: {_id:"$title",
            first: {$min:"$timestamp"},
            last:  {$max:"$timestamp"} }},
  {$project: {_id: 0,
              title: "$_id",                ($last-$first)/86400000
              age: {$divide:
                    [{$subtract: ["$last", "$first"]},
                     86400000]}},
              age_unit: {$literal: "day"}}}
])
```


Dealing with data of array type

- To aggregate (e.g. grouping) values in an array field, it is possible to flatten the array to access individual value
- **\$unwind** stage flattens an array field from the input documents to output a document for *each* element. Each output document is the input document with the value of the array field replaced by the element.
 - ▶ `{ $unwind: <field path> }` or
 - ▶

```
{
  $unwind:
    {
      path: <field path>,
      includeArrayIndex: <string>,
      preserveNullAndEmptyArrays: <boolean>
    }
}
```

\$unwind example

■ Default behaviour

▶ Input document:

```
{ "_id" : 1, "item" : "ABC1", "sizes" : [ "S", "M", "L" ] }
```

▶ After \$unwind: "\$sizes"

▶ Becomes 3 output documents:

```
{ "_id" : 1, "item" : "ABC1", "sizes" : "S" }
```

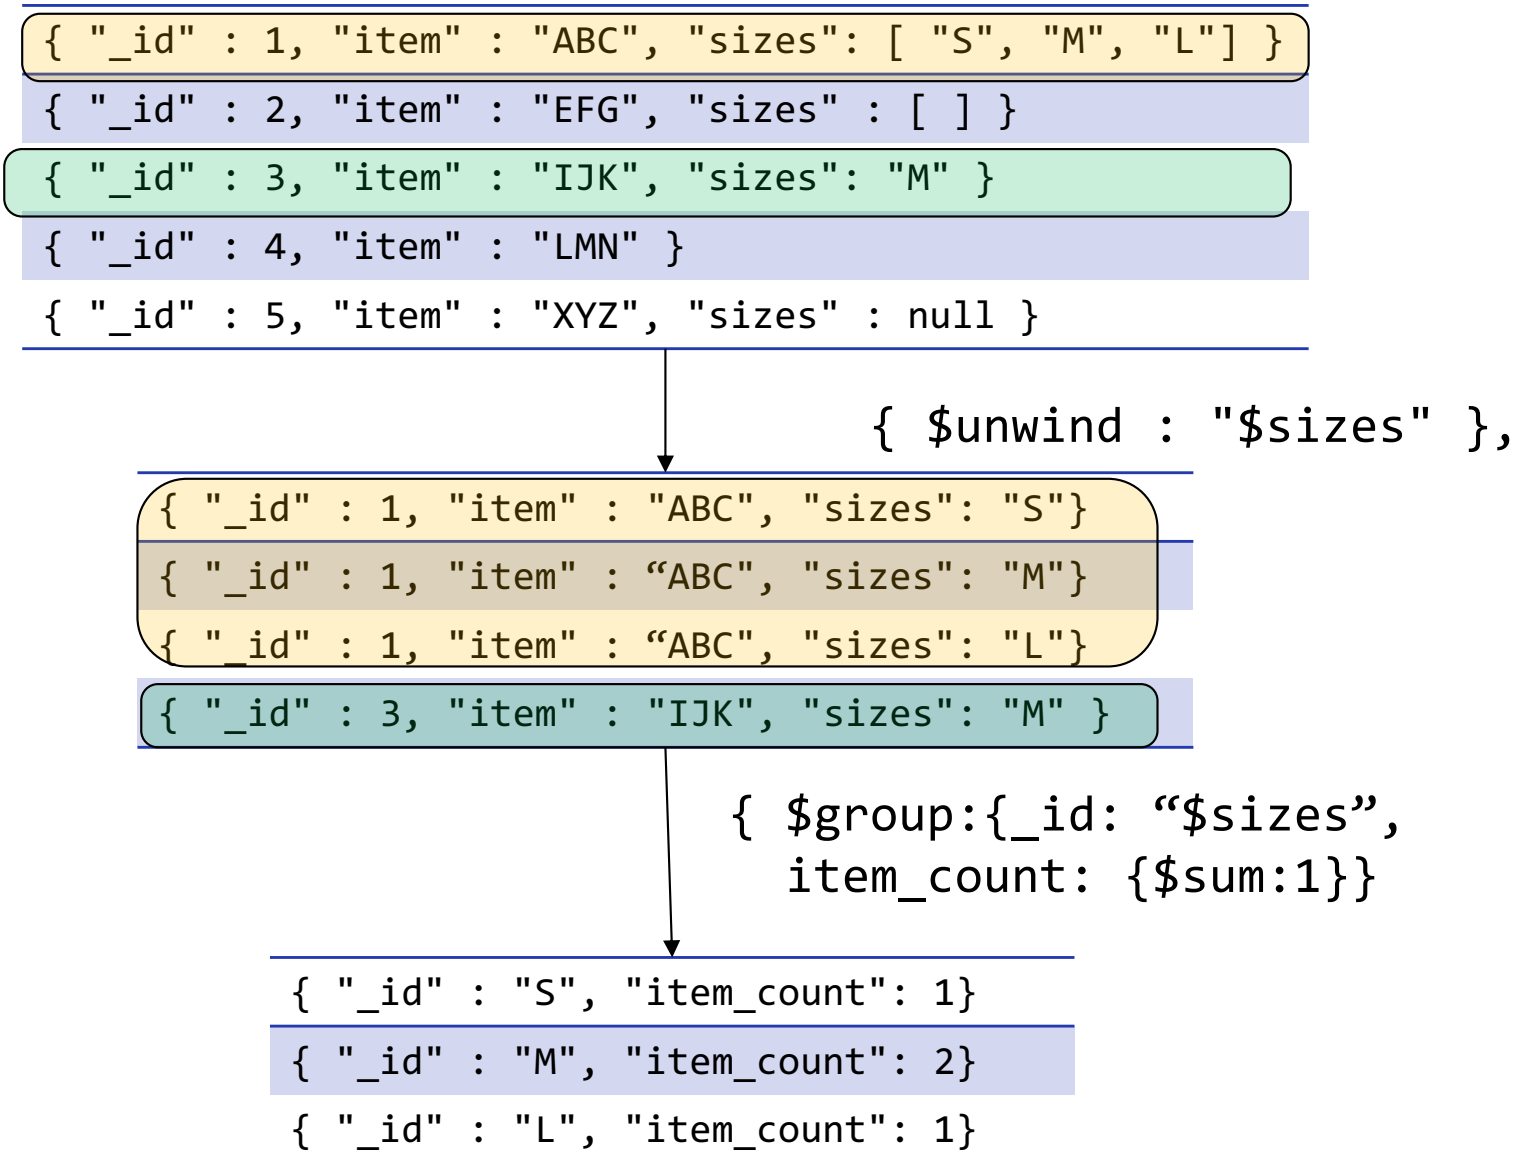
```
{ "_id" : 1, "item" : "ABC1", "sizes" : "M" }
```

```
{ "_id" : 1, "item" : "ABC1", "sizes" : "L" }
```

■ Find the number of items that are available in each size

```
db.inventory.aggregate( [  
    { $unwind : "$sizes" },  
    { $group: {_id: "$sizes", item_count: {$sum:1}} }  
] )
```

\$unwind then \$group



`$sort`, `$skip`, `$limit` and `$count` stages

- **`$sort`** stage sorts the incoming documents based on specified field(s) in ascending or descending order
 - ▶ The function and format is similar to the sort modifier in **`find`** query
 - ▶ `{ $sort: { <field1>: <sort order>, <field2>: <sort order> ... } }`
- **`$skip`** stage skips over given number of documents
 - ▶ The function and format is similar to the skip modifier in **`find`** query
 - ▶ `{ $skip: <positive integer> }`
- **`$limit`** stage limits the number of documents passed to the next stage
 - ▶ The function and format is similar to the limit modifier in **`find`** query
 - ▶ `{ $limit: <positive integer> }`
- **`$count`** stage counts the number of documents passing to this stage
 - ▶ The function and format is similar to the count modifier in **`find`** query
 - ▶ `{ $count: <string> }`
 - ▶ String is the name of the field representing the count

`$sample` and `$out` stages

- The **`$sample`** stage randomly selects given number of documents from the previous stage
 - ▶ `{ $sample: { size: <positive integer> } }`
 - ▶ Different sampling approaches depending on the location of the stage and the size of the sample and the collection
 - ▶ May fail due to memory constraints
- The **`$out`** stage writes the documents in a given collection
 - ▶ should be the last one in the pipeline
 - ▶ `{ $out: "<output-collection>" }`

Aggregation Operators

- A few aggregation stages allow us to add new fields or to give existing fields new values based on expression
 - ▶ In **\$group** stage we can use various *operators* or *accumulators* to compute values for new fields
 - ▶ In **\$project** stage we can use operators to compute values for new or exiting fields
- There are many predefined operators for various data types to carry out common operations in that data type
 - ▶ Arithmetic operators: **\$mod**, **\$log**, **\$sqrt**, **\$subtract**, ...
 - ▶ String operators: **\$concat**, **\$split**, **\$indexofBytes**, ...
 - ▶ Comparison operators: **\$gt**, **\$gte**, **\$lt**, **\$lte**,...
 - ▶ Set operators: **\$setEquals**, **\$setIntersection**, ...
 - ▶ Boolean operators: **\$and**, **\$or**, **\$not**, ...
 - ▶ Array operators: **\$in**, **\$size**, ..

Aggregation vs. Query operators

- There is another set of operators that can be used in **find/update/delete** queries or the **\$match** stage of an aggregation
 - ▶ E.g. **\$gt**, **\$lt**, **\$in**, **\$all**...
- The set is smaller and are different to the operators used in **\$group** or **\$project** stage
- Some operators look the same but have different syntax and slightly different interpretation in query and in aggregation.
 - ▶ E.g. **\$gt** in find query looks like
`{age: {$gt:18}}`
 - ▶ **\$gt** in **\$project** stage looks like:
`{over18: {$gt:["$age", 18]}}`

Returns true or false

Outline

■ Null type

■ Aggregation

- ▶ Single collection aggregation pipeline
- ▶ **Aggregation pipeline with multiple collections**

\$lookup stage

- **\$lookup** stage is added since 3.2 to perform left outer join between two collections
 - ▶ The collection already in the pipeline (maybe after a few stages)
 - ▶ Another collection (could be the same one)
- For each incoming document from the pipeline, the \$lookup stage adds a new **array field** whose elements are the matching documents from the other collection.
- A few different forms
 - ▶ Equality match
 - ▶ Join with other conditions
 - ▶ Join with uncorrelated sub-queries

\$lookup stage (cont'd)

- The output of **\$lookup** stage has the same number of documents as the previous stage
- Each document is augmented with an **array field** storing matching document(s) from the other collection
- The array could contain any number of documents depending on the match, including zero
- Missing local or foreign field is treated as having **null** value

Equality Match \$lookup

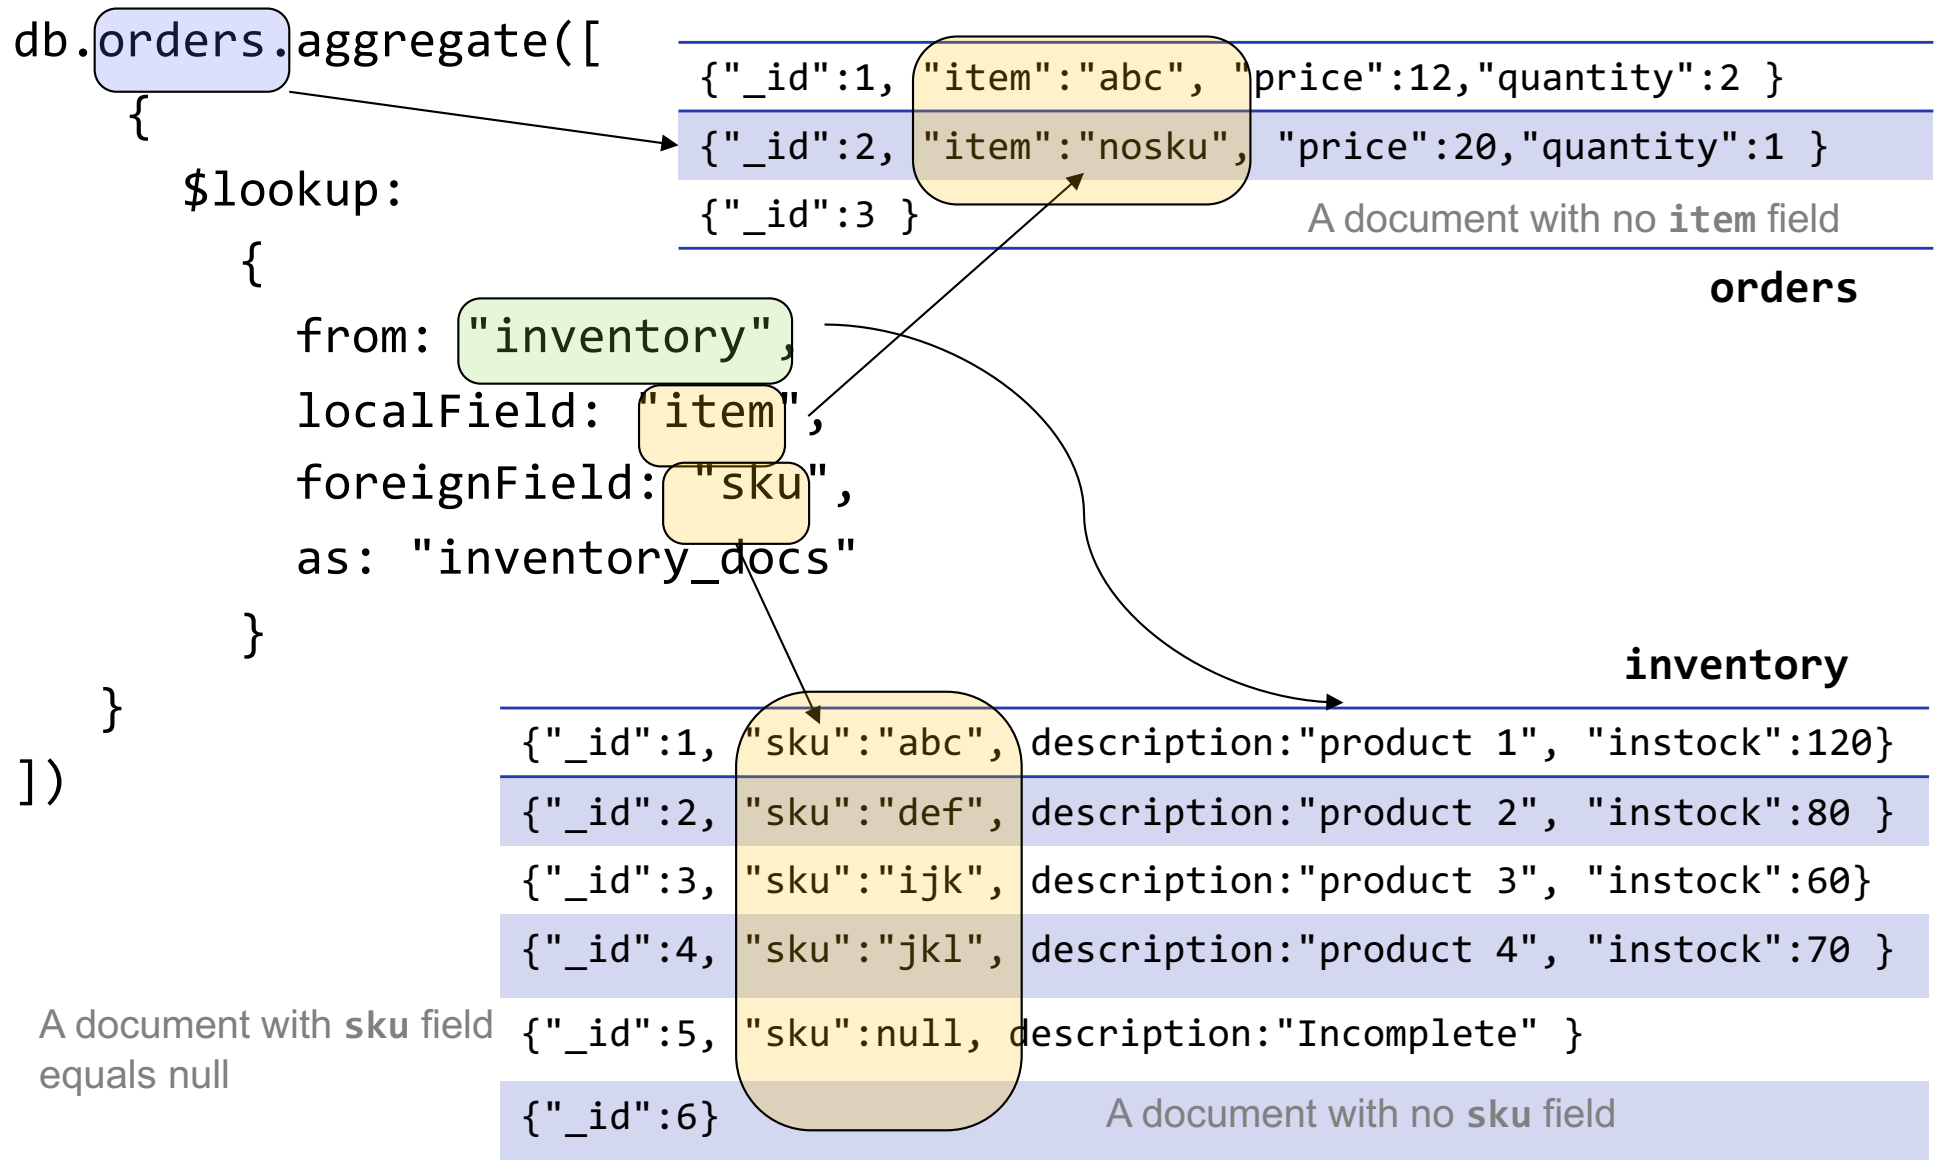
{ \$lookup:

```
{ from: <collection to join>,  
  localField: <field from the input documents>,  
  foreignField: <field from the documents of the "from" collection>,  
  as: <output array field>  
}
```

}

```
SELECT *, <output array field>  
FROM collection  
WHERE <output array field> IN (SELECT *  
                                FROM <collection to join>  
                                WHERE <foreignField>= <collection.localField>);
```

Equality match \$lookup example



https://docs.mongodb.com/manual/reference/operator/aggregation/lookup/#pipe._S_lookup

Equality match \$lookup example (cont'd)

```
{ "_id":1, "item":"abc", "price":12,"quantity":2 }
```

```
{ "_id":2, "item":"nosku", "price":20,"quantity":1 }
```

```
{ "_id":3 }
```

local

```
{ "_id":1, "sku":"abc", description:"product 1", "instock":120 }
```

```
{ "_id":2, "sku":"def", description:"product 2", "instock":80 }
```

```
{ "_id":3, "sku":"ijk", description:"product 3", "instock":60 }
```

```
{ "_id":4, "sku":"jkl", description:"product 4", "instock":70 }
```

```
{ "_id":5, "sku":null, description:"Incomplete" }
```

```
{ "_id":6 }
```

foreign

Non exists field matches null and non exists field

```
{ "_id":1, "item":"abc", "price":12,"quantity":2,
```

```
  "inventory_docs": [
```

```
    { "_id":1, "sku":"abc", description:"product 1", "instock":120 } ] }
```

```
{ "_id":2, "item":"nosku", "price":20,"quantity":1,
```

```
  "inventory_docs" : [ ] }
```

An empty array for no matching from other collection

```
{ "_id":3, "inventory_docs" : [
```

```
  { "_id" : 5, "sku" : null, "description" : "Incomplete" },
```

```
  { "_id" : 6 } ] }
```

output

Other format of \$lookup

```
{
  $lookup:
  {
    from: <collection to join>,
    let: { <var_1>: <expression>, ..., <var_n>: <expression> },
    pipeline: [ <pipeline to execute on the collection to join> ],
    as: <output array field>
  }
}
```

let: Optionally specifies variables to use in the pipeline field stages. Most likely the variable(s) may refer to field(s) in the local collection already in the pipeline

pipeline: Specifies the pipeline to run on the joined collection. The `pipeline` determines the resulting documents from the joined collection.

Multiple Joint Condition Example

orders collection

```
{ "_id" : 1, "item" : "almonds", "price" : 12, "ordered" : 2 }  
{ "_id" : 2, "item" : "pecans", "price" : 20, "ordered" : 1 }  
{ "_id" : 3, "item" : "cookies", "price" : 10, "ordered" : 60 }
```

warehouses collection

```
{ "_id" : 1, "stock_item" : "almonds", warehouse: "A", "instock" : 120 }  
{ "_id" : 2, "stock_item" : "pecans", warehouse: "A", "instock" : 80 }  
{ "_id" : 3, "stock_item" : "almonds", warehouse: "B", "instock" : 60 }  
{ "_id" : 4, "stock_item" : "cookies", warehouse: "B", "instock" : 40 }  
{ "_id" : 5, "stock_item" : "cookies", warehouse: "A", "instock" : 80 }
```

An ordered item may be stocked in multiple warehouses;

We want to find for each ordered item the warehouse with sufficient stock to cover the order

Multiple Joint Condition

- This query involves comparing two fields of the local and foreign documents:
 - ▶ “item” in **orders** should match “stock_item” in **warehouses**
 - ▶ “ordered” in **orders** should be less than or equal to “instock” in **warehouses**

orders collection	
{ "_id" : 1, "item" : "almonds", "price" : 12, "ordered" : 2 }	
{ "_id" : 2, "item" : "pecans", "price" : 20, "ordered" : 1 }	
{ "_id" : 3, "item" : "cookies", "price" : 10, "ordered" : 60 }	
warehouses collection	
{ "_id" : 1, "stock_item" : "almonds", warehouse: "A", "instock" : 120 }	
{ "_id" : 2, "stock_item" : "pecans", warehouse: "A", "instock" : 80 }	
{ "_id" : 3, "stock_item" : "almonds", warehouse: "B", "instock" : 60 }	
{ "_id" : 4, "stock_item" : "cookies", warehouse: "B", "instock" : 40 }	
{ "_id" : 5, "stock_item" : "cookies", warehouse: "A", "instock" : 80 }	

Multiple Joint Condition \$lookup

```
db.orders.aggregate([
  {
    $lookup:
    {
      from: "warehouses",
      let: { order_item: "$item", order_qty: "$ordered" },
      pipeline: [
        { $match:
          { $expr:
            { $and:
              [
                { $eq: [ "$stock_item", "$$order_item" ] },
                { $gte: [ "$instock", "$$order_qty" ] }
              ]
            }
          }
        }
      ],
      as: "stockdata"
    }
  }
])
```

This is the way to specify multiple condition

This is the way to let the pipeline access local fields: use variable **order_item** to access the local document's **item** field; use variable **order_qty** to access the local document's **ordered** field

variables are accessed using "\$\$" prefix

\$lookup by default includes the entire matched foreign document in the array, we can use \$project stage to get rid of some field

Matching document after the pipeline stage will be stored in this variable

Results

```
{ "_id" : 1,  
  "item" : "almonds",  
  "price" : 12,  
  "ordered" : 2,  
  "stockdata" : [  
    { "warehouse" : "A", "instock" : 120 },  
    { "warehouse" : "B", "instock" : 60 }  
  ]  
}
```

```
{ "_id" : 2,  
  "item" : "pecans",  
  "price" : 20,  
  "ordered" : 1,  
  "stockdata" : [  
    { "warehouse" : "A", "instock" : 80 }  
  ]  
}
```

```
{ "_id" : 3,  
  "item" : "cookies",  
  "price" : 10,  
  "ordered" : 60,  
  "stockdata" : [  
    { "warehouse" : "A", "instock" : 80 }  
  ]  
}
```

Uncorrelated Subquery Example

absences collection

```
{
  "_id" : 1,
  "student" : "Ann Aardvark",
  sickdays: [ "2018-05-01", 2018-08-23" ]
}
```

```
{
  "_id" : 2,
  "student" : "Zoe Zebra",
  sickdays: [ "2018-02-01", 2018-05-23" ]
}
```

holidays collection

```
{ "_id" : 1, year: 2018, name: "New Years", date: "2018-01-01" }
```

```
{ "_id" : 2, year: 2018, name: "Pi Day", date: "2018-03-14" }
```

```
{ "_id" : 3, year: 2018, name: "Ice Cream Day", date: "2018-07-15" }
```

```
{ "_id" : 4, year: 2017, name: "New Years", date: "2017-01-01" }
```

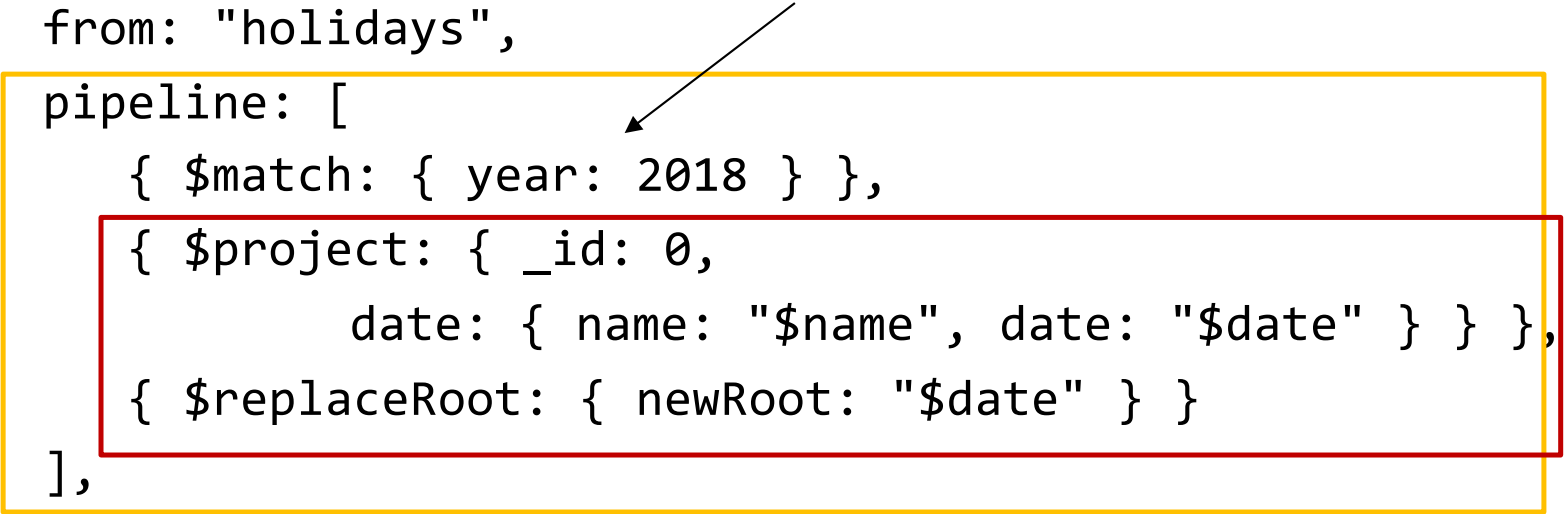
```
{ "_id" : 5, year: 2017, name: "Ice Cream Day", date: "2017-07-16" }
```

We want to include all 2018 public holidays in the **absences** collection

Uncorrelated Subquery \$lookup

```
db.absences.aggregate([
  {
    $lookup:
      {
        from: "holidays",
        pipeline: [
          { $match: { year: 2018 } },
          { $project: { _id: 0,
            date: { name: "$name", date: "$date" } } },
          { $replaceRoot: { newRoot: "$date" } }
        ],
        as: "holidays"
      }
  }
])
```

The inner pipeline selects documents from **holidays** collection based on a condition unrelated with the local collection



```
    { $match: { year: 2018 } },
    { $project: { _id: 0,
      date: { name: "$name", date: "$date" } } },
    { $replaceRoot: { newRoot: "$date" } }
```

The **\$project** and **\$replaceRoot** change the structure of the inner pipeline output documents

Has the same effect as: **{ \$project: { _id: 0, year:0 } }** in this case

\$replaceRoot and similar “project” like stage are useful for promoting an embedded document at root level

Results

```
{
  "_id" : 1,
  "student" : "Ann Aardvark",
  "sickdays" : [ "2018-05-01", "2018-08-23"],
  "holidays" : [
    { "name" : "New Years", "date" : "2018-01-01" },
    { "name" : "Pi Day", "date" : "2018-03-14" },
    { "name" : "Ice Cream Day", "date" : "2018-07-15" }
  ]
}
```

```
{
  "_id" : 2,
  "student" : "Zoe Zebra",
  "sickdays" : [ "2018-02-01", "2018-05-23" ],
  "holidays" : [
    { "name" : "New Years", "date" : "2018-01-01" },
    { "name" : "Pi Day", "date" : "2018-03-14" },
    { "name" : "Ice Cream Day", "date" : "2018-07-15" }
  ]
}
```

References

■ BSON types

- ▶ <https://docs.mongodb.com/manual/reference/bson-types/>

■ Querying for Null or Missing Field

- ▶ <https://docs.mongodb.com/manual/tutorial/query-for-null-fields/index.html>

■ Aggregation Pipelines

- ▶ <https://docs.mongodb.com/manual/core/aggregation-pipeline/>

■ Aggregation operators

- ▶ <https://docs.mongodb.com/manual/reference/operator/aggregation/>