

COMP5338 – Advanced Data Models

Week 7: Neo4j Internal and Data Modelling

Dr. Ying Zhou
School of Computer Science



Property Graph Model




Table concept is not part of the data model

The database is a large graph, could contain independent subgraphs

Community edition only supports one user database; Enterprise edition supports multiple user databases

Logic data model and physical storage model could be totally different

It is theoretically possible to construct a property graph model with any storage backend

Outline

■ Neo4j Storage

■ Neo4j Query Plan and Indexing

■ Neo4j – Data Modeling

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1968

WARNING

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (the Act).

The material in this communication may be subject to copyright under the Act. Any further reproduction or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice

Materials adapted by permission from *Graph Databases (2nd Edition)* by Ian Robinson et al (O'Reilly Media Inc.). Copyright 2015 Neo Technology, Inc

COMP5338 "Advanced Data Models" - 2020 (Y. Zhou)

07-2

Index-free Adjacency

■ Native storage of relationships between nodes

- ▶ Effectively a pre-computed bidirectional join

■ Traversal is like pointer dereferencing


- ▶ Almost as fast as well

■ Index-free Adjacency

- ▶ Each node maintains a direct link to its adjacent nodes
- ▶ Each node is effectively a micro-index to the adjacent nodes

■ Cheaper than global indexes

- ▶ Queries are faster, do not depend on the total size of the graph




Slides 4-11 are based on Graph Database chapter 6.1 and 6.2

COMP5338 "Advanced Data Models" - 2020 (Y. Zhou)

07-4

Neo4j Architecture




Page 163 of Graph Databas



Node store file

- All node data is stored in **one** node store file
 - Physically stored in file named *neostore.nodestore.db*
 - Each record is of a **fixed size** – 15 bytes (*was 9 bytes in earlier version*)
 - Offset of stored node = node id * 15 (node id = 100, offset = 1500)
 - Deleted IDs in *.id* file and can be reused



Property Graph and Store files


- Graph data is stored in *store files* on disk
 - ▶ Nodes, relationships, properties, labels and types all have their own store files.
 - Check under <neo4j-home>/data/databases/neo4j
 - ▶ Separating graph structure and property data promotes fast traversal
 - Node, relationship, property, label and type all have system assigned IDs
 - They are stored as fixed length record in respective stores
 - user's view of their graph and the actual records on disk are structurally dissimilar

—	use Neo4j's view of their graph and the details of which nodes are structurally dissimilar
48K	1 Oct 22:56 neostore.relationshipgroupstore.db. ub.id



Relationship store file

- All relationship data is stored in **one** relationship store file
 - Physically stored in file named *neostore.relationshipstore.db*
 - Each record is of a fixed size – 34 bytes
 - Offset of stored relationship = relationship id * 34
 - ▶ So, relationship id = 10, offset = 340



Implications


Both Node ID and Property ID are of 4 bytes

- The maximum ID value is $2^{32} - 1$
 - There is a maximum number of nodes/relationships in a database
- ID is assigned and managed by the system
 - The corresponding record will be stored in the computed offset
- The IDs of deleted nodes/relationships will be reused

07-9

Node structure

Bob LIKES Alice



07-11

Other Files

Property store contains fixed size records to store properties for nodes and relationships

- Simple properties are stored inline
- Complex ones such as long string or array property are stored elsewhere


Node label in node records references data in label store

Relationship type in relationship record references data in relationship type store


10K	1 Oct 22:56	neostore.labels.db
568K	1 Oct 22:56	neostore.nodestore.db.id
88K	1 Oct 22:56	neostore.nodestore.db.labels
8.0K	1 Oct 21:09	neostore.nodestore.db.labels
40K	1 Oct 22:56	neostore.nodestore.db.labels.id
48K	1 Oct 22:56	neostore.propertystore.db
1.1M	1 Oct 22:56	neostore.relationshipstore.db
96K	1 Oct 22:56	neostore.relationshipstore.db.id
8.0K	1 Oct 22:17	neostore.relationshiptypestore.db
40K	1 Oct 22:56	neostore.relationshiptypestore.db.id
8.0K	1 Oct 22:17	neostore.relationshiptypestore.db.names
40K	1 Oct 22:56	neostore.relationshiptypestore.db.names
48K	1 Oct 22:56	neostore.propertystore.db.strings
48K	1 Oct 22:56	neostore.propertystore.db.strings.id

Relationship structure

Bob LIKES Alice



Doubly linked list




*"The node and relationship stores are concerned **only** with the **structure** of the graph, not its property data. Both stores use fixed-sized records so that any individual record's location within a store file can be rapidly computed given its ID. These are critical design decisions that underline Neo4j's commitment to high-performance traversals."*

-- Chapter 6, Graph Databases


Doubly linked list (cont'd)

Add another node and relationship
create node Daniel
relationship Charlie --> Daniel

Node file



Relationship file



Outline

■ Neo4j Storage

■ Neo4j Query Plan and Indexing

■ Neo4j – Data Modeling



Neo4j Query Execution

- Each Neo4j Query is turned into an execution plan by an **execution planner**
- The **execution plan** is a tree-like structure consists of various **operators**, each implements a specific piece of work
- Query plan stages
 - Starting point (leaf node)
 - Obtaining data from storage engine
 - Expansion by matching given pattern in the query statement
 - Row filtering, skipping, sorting, projection, etc...
 - Combining operations
 - Updating
- Execution plans are evaluated based on statistics maintained by database
 - The number of nodes having a certain label.
 - The number of relationships by type.
 - Selectivity per index.
 - The number of relationships by type, ending with or starting from a node with a specific label.




Evaluation Statistics

- Each operator is annotated with some statistics
- Rows:** The number of rows that the operator produced. This is only available if the query was *profiled*.
- EstimatedRows:** This is the estimated number of rows that is expected to be produced by the operator.
- DbHits:** Some operator needs to retrieve data from or update data in the storage. A *database hit* is an abstract unit of this storage engine work
 - Creating a node, a relationship, a label, a type
 - Deleting a node, a relationship,
 - Getting a node, a property of a node, the label,...
 - Getting a relationship, a property of a label, the type,
 - Updating...



Query Plan: an example



This is the performance output from Neo4j 3.x

Query Starting Points



- Most queries start with one or a set of **nodes** except if a relationship ID is specified
 - MATCH (n1)-[r]->() WHERE id(r)=0 RETURN r, n1**
 - This query will start from locating the first record in the relationship file
- Query may start by scanning all nodes
 - MATCH(n) RETURN (n)**
 - MATCH (cloudAtlas {title: "Cloud Atlas"})<[:DIRECTED]-{directors} RETURN directors.name**
- Query may start by scanning all nodes belonging to a given label
 - MATCH (p:Person{name:"Tom Hanks"}) return p**
 - Labels are implicitly indexed
- Query may start by using index



Query starting from labelled node

```

MATCH (n:Person) -[r]-(something)
WITH n, count(something) as degree
ORDER BY degree DESC
LIMIT 1
RETURN n, degree
  
```




A relatively complex query and plan

```


MATCH (n:Person{name: "Tom Hanks"}) ←
WITH n.phone as phones, n
UNWIND phones as phone
MATCH (m:Person) ←
WHERE phone in m.phone and n<>m
RETURN m.name
  
```

Apply works by performing a nested loop. Every row being produced on the left-hand side of the **Apply** operator will be fed to the leaf operator on the right-hand side, and then **Apply** will yield the combined results



Query Plan With Index

- Neo4j supports index on properties of labelled node
- Index has similar behaviour as those in relational systems
- It can be built on single or composite properties
- Create Index
 - CREATE INDEX ON :Person(name)**
- Drop Index
 - DROP INDEX ON :Person(name)**



Each row represents a path that may include more than one relationship, so the db hits number is larger than row number

Query:



```

MATCH [bacon:Person {name:"Kevin Bacon"}]-[*1..4]-(hollywood)
RETURN DISTINCT hollywood
  
```

Comparing Execution Plans

```

MATCH (m:Movie)
WHERE m.released IN [1999,2003]
RETURN m.title, m.released
  
```






For each year value, all Movie nodes are examined. The entire Movie nodes set is retrieved and examined twice

Comparing Execution Plans (with Index)

```

MATCH (m: Movie)
WHERE m.released IN [1999,2003]
RETURN m.title, m.released
    
```




Similar performance

```

UNWIND [1999,2003] as year
MATCH (m: Movie)
WHERE m.released = year
RETURN m.title, m.released
    
```



For each year value, an index search is performed



CREATE INDEX ON :Movie(released)

COMP5338 "Advanced Data Models" - 2020 (Y. Zhou)

07-25

Option 1


```

MATCH (p:Person)-[:ACTED_IN]->(m:Movie)
WITH p, count(m) AS mc
WHERE mc > 2
MATCH (p)-[:DIRECTED]->(m2:Movie)
RETURN p.name, m2.title
    
```

Start with Movie label because there are less Movie nodes than Person nodes



15 person have acted in at least 3 movies



Need to check all 15 person's relationships


Check the other node is movie node

Getting two properties

Aggregate by p

Check the other nodes are of Person type

Aggregate by p



07-27

Another example of comparison

■ Question: Find out a list of person who has acted in at least three movies and also directed at least one movie

■ Cypher is powerful and flexible

- ▶ It is possible to write very different queries that produce the same results
- ▶ The performance could have big difference
- ▶ The DB engine does not have much knowledge to rewrite the queries as those in SQL
 - Not based on relational algebra

COMP5338 "Advanced Data Models" - 2020 (Y. Zhou)


07-26

Option 2

```

MATCH (m1:Movie)<-[a:ACTED_IN]-(p:Person)-[:DIRECTED]->(m2:Movie)
WITH p, count(distinct m1) as ac, m2
WHERE ac > 2
RETURN p.name, m2.title
    
```

Obtain the 18 Movie nodes from the 18 relationships as m1



38 Movie nodes + 44 relationships

44 Person nodes

Among 62 relationships only 18 are of ACTED type

38 estimated rows 39 db hits

44 estimated rows 82 db hits

44 estimated rows 44 db hits

44 estimated rows 44 db hits

57 estimated rows 62 db hits

Aggregation is processed in memory.

Filtering

57 estimated rows 18 db hits

8 estimated rows 6 db hits

2 estimated rows 2 db hits

5.81 total memory (bytes) 2 db hits

COMP5338 "Advanced Data Models" - 2020 (Y. Zhou)

07-28

Performance of Creation Operations

- Example: add COSTARTS relationships for Tom Hanks
- Option 1: using CREATE

```

1 profile
2 MATCH (hanks:Person {name:"Tom Hanks"})
3   -[:ACTED_IN]→(m)←[:ACTED_IN]-(p:Person)
4 Where hanks↔ p
5 WITH hanks, collect(distinct p) as costars
6 unwind costars as p
7 create (hanks)-[:COSTARS]→(p)

```



Option 2: Using MERGE


```

1 profile
2 MATCH (hanks:Person {name:"Tom Hanks"})
3   -[:ACTED_IN]→(m)←[:ACTED_IN]-(p:Person)
4 Where hanks↔ p
5 WITH hanks, p
6 merge (hanks)-[:COSTARS]→(p)

```



Profile Result



Profile Result

For each p , check if there is a COSTARS relationship between node $hanks$ and node p .


Check relationship existence needs to traverse the entire relationship link.

Create the relationship when it does not exists. There are 34 create relationship actions.

```

1 profile
2 MATCH (hanks:Person {name:"Tom Hanks"})
3   -[:ACTED_IN]→(m)←[:ACTED_IN]-(p:Person)
4 Where hanks↔ p
5 WITH hanks, p
6 merge (hanks)-[:COSTARS]→(p)

```



Transactions

- Neo4j supports full ACID transactions
 - ▶ Similar to those in RDBMS
- Uses locking to ensure consistency
 - ▶ Lock Manager manages locks held by a transaction
- Logging
 - ▶ Write Ahead Logging (WAL)
- Transaction Commit Protocol
 - ▶ Acquire locks (Atomicity, Consistency, Isolation)
 - ▶ Write Undo and Redo records to the WAL
 - for each node, relationship, property changed is written to the log
 - ▶ Write commit record to the log and flush to disk (Durability)
 - ▶ Release locks
- Recovery – if the database server/machine crashes
 - ▶ Apply log records to replay changes made by the transactions



Graph Data Modelling

- Graph data modelling is very closely related with domain modelling
- You need to decide
 - ▶ Node or Relationship
 - ▶ Node or Property
 - ▶ Label/Type or Property
- Decisions are based on
 - ▶ Features of entities in application domain
 - ▶ Your typical queries
 - ▶ Features and constraints of the underlying storage system




Outline

- Neo4j Storage
- Neo4j Query Plan and Indexing
- Neo4j – Data Modeling



Node vs. Relationship

- Nodes for Things, Relationship for Structures
 - ▶ AS A reader who likes a book, I WANT to know which books other readers who like the same book have liked, SO THAT I can find other books to read.




```
MATCH (:Reader {name:'Alice'})-[:LIKES]->(:Book {title:'Dune'})  
<-[:LIKES]-(:Reader)-[:LIKES]->(books:Book)  
RETURN books.title
```



Node vs. Relationship

■ Model Facts as Nodes




Relationship Property or Relationship Type

- E.g. The relationship between user node and address node can be:
 - typed as **HOME_ADDRESS**, **BILLING_ADDRESS** or
 - typed as generic **ADDRESS** and differentiated using a type property {type:'home'}, {type:'billing'}
- We use fine-grained relationships whenever we have a closed set of relationship types.
 - ▶ Eg. there are only a finite set of address types
 - ▶ If traversal would like to follow generic type **ADDRESS**, we may have to use redundant relationships
 - MATCH (user)-[:**HOME_ADDRESS**|**WORK_ADDRESS**|**DELIVERY_ADDRESS**]->(address)
 - MATCH (user)-[:**ADDRESS**]->(address)
 - MATCH (user:User)-[r]->(address:Address)



Node or Property

■ Represent Complex Value Types as Nodes



References

- Ian Robinson, Jim Webber and Emil Eifrem, *Graph Databases*, Second Edition, O'Reilly Media Inc.,
 - ▶ You can download this book from the Neo4j site, <https://neo4j.com/graph-databases-book/?ref=home>
 - Chapter 4, Chapter 6
- Neo4j – Reference Manual
 - ▶ <https://neo4j.com/docs/developer-manual/current/>
- Neo4j reference manual: Execution plan operators in detail
 - ▶ <https://neo4j.com/docs/cypher-manual/current/execution-plans/operators/>

