

COMP5338 – Advanced Data Models

Week 9: Spatial Data Model and Query

Dr. Ying Zhou

School of Computer Science



THE UNIVERSITY OF
SYDNEY

Outline

■ Motivation

■ Spatial Data Model

■ Spatial Data Queries

■ Spatial Query in MongoDB



Motivation

- Many entities represent physical objects, and some physical features also matter for business purposes
 - ▶ Eg. A store has a name, business category, contact number and is located at a particular place (geo-spatial)
 - The geo-spatial feature can help find a store that is close to a customer
 - ▶ Eg. A toy has name, category, material, and shape
 - The shape feature helps to find a box that can fit the toy
- There are large amount of geospatial data:
 - ▶ Businesses and homes have addresses
 - Both the logic aspect and physical aspect
 - ▶ Google Maps, Google Earth
 - ▶ Weather and Climate Data



Spatial Data and Object Concept

- Spatial feature could refer to a
 - ▶ Point, a 2d shape, a 3d shape, or shape in higher dimension
- Spatial features of an entity cannot be represented as simple value type
 - ▶ A point in 2d space has two coordinate values
- It is natural to use object to represent spatial features (spatial data)
- There are also spatial related operations we need to perform on spatial data
 - ▶ Compute distance between points
 - ▶ Compute area of 2d shape or volume of 3d shapes
 - ▶ Compute various relationships among spatial objects

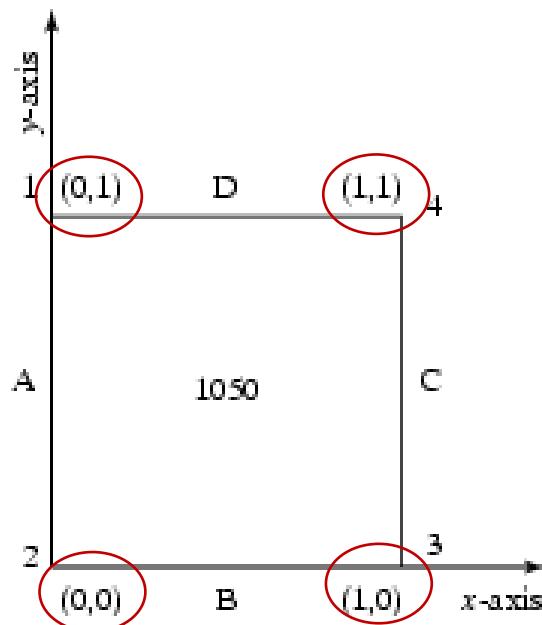


Spatial Data as Object Example

- Consider a spatial data representing census block:

```
CREATE TABLE census_blocks (
    name      string,
    area      float,
    population   number,
    boundary    Polyline );
```

A User
Defined Type



Census_blocks			
Name	Area	Population	Boundary
1050	1	1839	Polyline((0,0),(0,1),(1,1),(1,0))

From Shekhar-Chawla, Fig 1.4

Spatial Data in Purely Relational Form

Census_blocks

Name	Area	Population	boundary-ID
340	1	1839	1050

Polygon

boundary-ID	edge-name
1050	A
1050	B
1050	C
1050	D

Edge

edge-name	endpoint
A	1
A	2
B	2
B	3
C	3
C	4
D	4
D	1

Point

endpoint	x-coor	y-coor
1	0	1
2	0	0
3	1	0
4	1	1

From Shekhar-Chawla, Fig 1.4



Spatial Database Management System

- A SDBMS is a software module that
 - ▶ can work with an underlying DBMS
 - ▶ supports spatial data models, spatial abstract data types (ADTs) and a query language from which these ADTs are callable
 - ▶ supports spatial indexing, efficient algorithms for processing spatial operations, and domain specific rules for query optimization
 - ▶ Many RDBMS and NoSQL storage systems have support for spatial data
 - Oracle, SQL Server, MongoDB, Neo4j, ...
- SDBMS components
 - ▶ spatial data model
 - ▶ query language
 - ▶ query processing
 - ▶ file organization and indices
 - ▶ query optimization
 - ▶ etc.



Outline

■ Motivation

■ Spatial Data Model

- ▶ Field vs. Object Models
- ▶ Coordinate System
- ▶ Topological Operations

■ Spatial Data Queries

■ Spatial Query in MongoDB



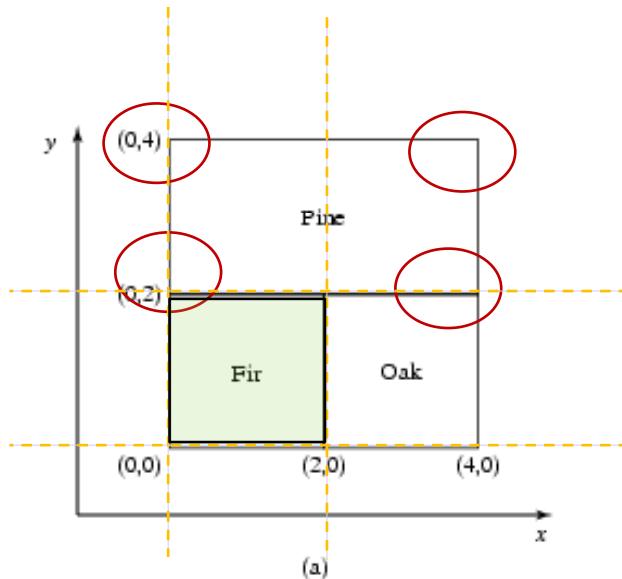
Models of Spatial Information

■ Two common models

- ▶ Field based (*also: space-based*)
 - Model properties of underlying space
 - Good for expressing values vary continuously over space (e.g. temperature, rainfall, elevation, depth, etc.)
 - Fields are actually **functions** that map spatial locations to values
- ▶ Object based (*also: entity-based*)
 - Model boundaries of spatial feature (e.g. the census block is modelled by a polygon)
 - Good for expressing discrete spatial entities



Examples of Field and Object Models



- (a) forest stand map
- (b) Object model has 3 polygons defining the boundaries
- (c) Field model uses a function to calculate the property "tree species"

Object Viewpoint of Forest Stands

Area-ID	Dominant Tree Species	Area/Boundary
FS1	Pine	$[(0,2),(4,2),(4,4),(0,4)]$
FS2	Fir	$[(0,0),(2,0),(2,2),(0,2)]$
FS3	Oak	$[(2,0),(4,0),(4,2),(2,2)]$

(b)

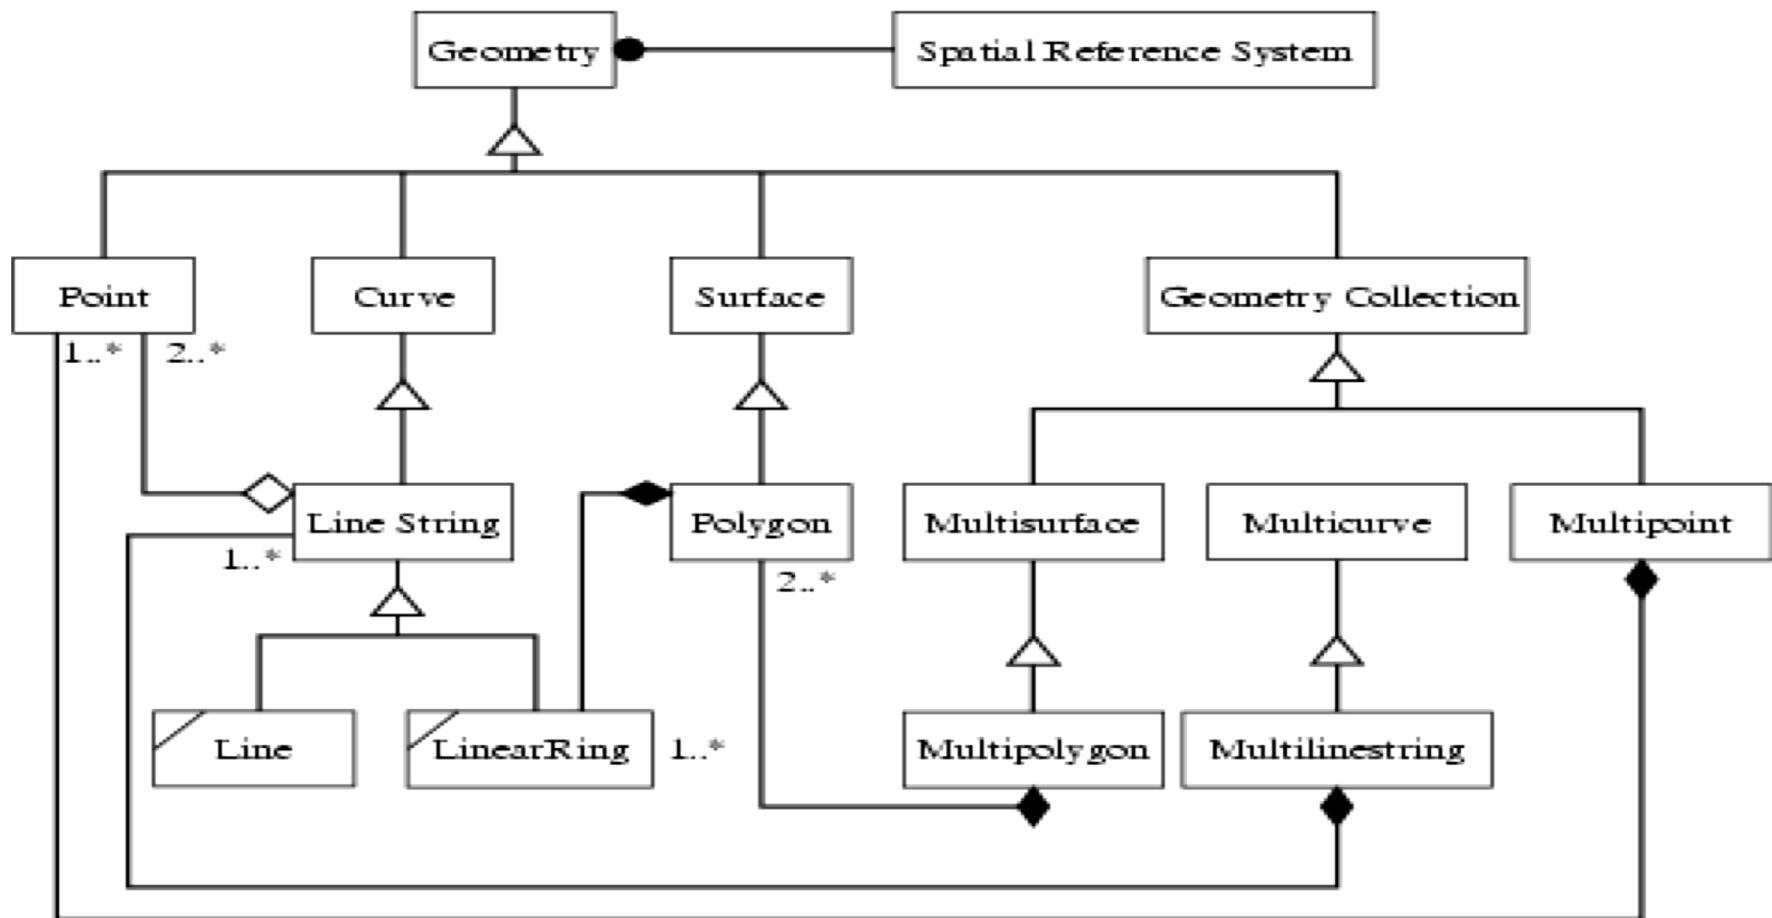
Field Viewpoint of Forest Stands

$$f(x,y) = \begin{cases} \text{"Pine," } 2 \leq x \leq 4; 2 < y \leq 4 \\ \text{"Fir," } 0 \leq x \leq 2; 0 \leq y \leq 2 \\ \text{"Oak," } 2 < x \leq 4; 0 \leq y \leq 2 \end{cases}$$

(c)

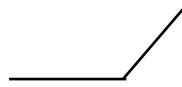
From Shekhar-Chawla, Fig 2.1

OpenGIS Geometry Model



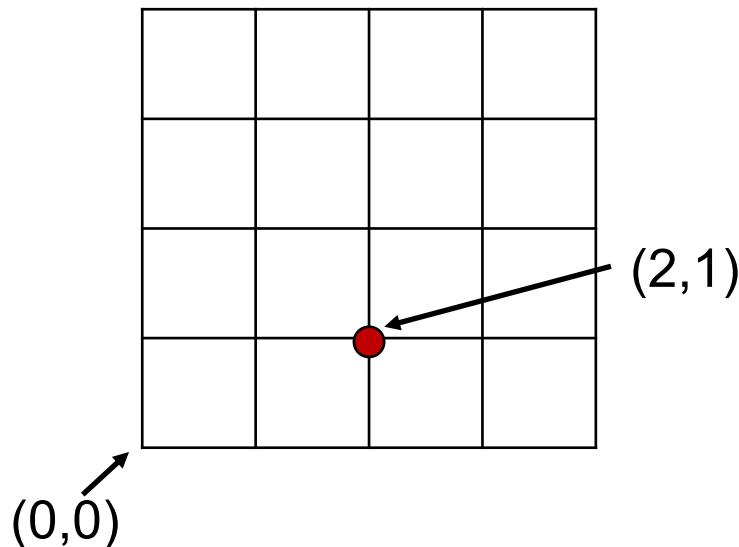
From Shekhar-Chawla, Fig 2.2

Example Spatial Objects

- **Point:** represented by its coordinates eg (-10, 30)
- Collection of several **points**
- **Line String**
 - ▶ Simplest form is piece-wise linear, given by points (and implying the straight segments between them)
 - ▶ Eg (0,1), (1,1), (2,2)
- **Polygon**
 - ▶ A 2-d region whose boundary is given
 - ▶ Simplest form: boundary is a **line string** that returns to its start
 - ▶ More complicated: region with holes
- Collection of **polygons**

Coordinate Systems

- Points from a 2-d space are represented by pairs of numbers
 - ▶ The numbers could refer to a dot on a drawing area, a piece of land in a game setting, or a location on earth
- There are many ways to associate numbers with points
- Simple 2d Cartesian coordinate
 - ▶ Choose a point as origin $(0,0)$
 - ▶ Choose a direction for the x-axis, and a scale (how far is $(1,0)$) from $(0,0)$?)



A Round World

- The surface of the earth is 2-dimensional, but curved
 - ▶ Cartesian systems work reasonably in small regions
- Traditional geographic coordinate system
 - ▶ 2d: longitude and latitude
 - ▶ 3d: longitude, latitude, elevation
 - ▶ The surface is a sphere
 - (-179, 10) is very close to (179, 10)
 - A linestring might cross the (long = 180) line;
- Many SDBMS supports both flat space and sphere



Operations on Spatial Objects in the Object Model

■ Classifying operations

- ▶ ***Set based:***
 - a set operation (e.g. intersection) of 2 polygons produce another polygon
- ▶ ***Topological operations:*** Boundary of USA touches boundary of Canada
- ▶ ***Directional:*** New York city is to east of Chicago
- ▶ ***Metric:*** Chicago is about 700 miles from New York city.

Set theory based	Union, Intersection, Containment
Topological	Touches, Disjoint, Overlap, etc.
Directional	East, North-West, etc.
Metric	Distance



Topological Relationships

■ Topological Relationships

- ▶ invariant under elastic deformation (without tear, merge).
- ▶ Two countries which touch each other in a planar paper map will continue to do so in spherical globe maps.

■ Example queries with topological operations

- ▶ What is the topological relationship between two objects A and B ?
- ▶ Find all objects which have a given topological relationship to object A?
 - E.g. find all rivers that cross a city

■ Can we express topological relationship mathematically?

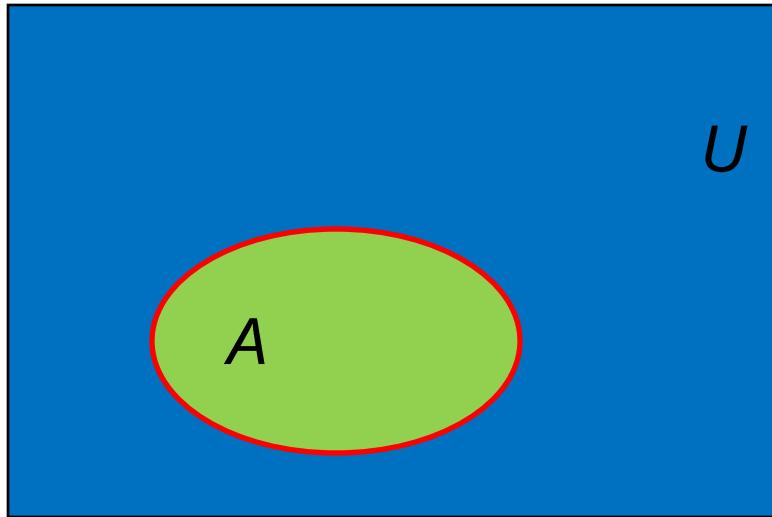
- ▶ Metric operations may be expressed using various functions, e.g. distance function
- ▶ Set operations can be express mathematically
- ▶ The mathematical form helps to calculate such relationships



Topological Concepts

■ Interior, boundary, exterior

- ▶ Let A be an object in a “Universe” U .



Green is A interior (A^o)

Red is boundary of A (∂A)

Blue – (Green + Red) is
 A exterior (A^-)

- ▶ Exterior is also referred to as the *complement* of an object

Nine-Intersection Model of Topological Relationships

- Many topological Relationship between A and B can be specified using 9 intersection model
 - ▶ Eight possible 2D topological relationships for objects without holes;
- Nine intersections
 - ▶ intersections between interior, boundary, exterior of A, B
 - ▶ A and B are spatial objects in a two dimensional plane.
 - ▶ Can be arranged as a 3 by 3 matrix
 - ▶ Matrix element take a value of 0 (false) or 1 (true)

$$\Gamma_9(A, B) = \begin{pmatrix} A^0 \cap B^0 & A^0 \cap \partial B & A^0 \cap B^- \\ \partial A \cap B^0 & \partial A \cap \partial B & \partial A \cap B^- \\ A^- \cap B^0 & A^- \cap \partial B & A^- \cap B^- \end{pmatrix}$$

From Shekhar-Chawla, p 28



Specifying Topological Operations using the 9-Intersection Model

	I	B	E
I			
B			
E			

A

B

For **disjoint** relation:
A's exterior intersects with
B's everything and vice
versa

If A **contains** B:
A's interior intersects with
B's everything and B's
exterior intersects with A's
everything

 $\begin{pmatrix} 0 & 0 & 1 \\ 0 & 0 & 1 \\ 1 & 1 & 1 \end{pmatrix}$ disjoint	 $\begin{pmatrix} 1 & 1 & 1 \\ 0 & 0 & 1 \\ 0 & 0 & 1 \end{pmatrix}$ contains	 $\begin{pmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 1 & 1 & 1 \end{pmatrix}$ inside	 $\begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$ equal
 $\begin{pmatrix} 0 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}$ meet	 $\begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{pmatrix}$ covers	 $\begin{pmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 1 & 1 & 1 \end{pmatrix}$ coveredBy	 $\begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}$ overlap

From Shekhar-Chawla, Fig 2.3



Outline

■ Motivation

■ Spatial Data Model concepts

■ Spatial Data Queries

- ▶ Query type
- ▶ General processing steps

■ Spatial Query in MongoDB



Spatial Processing

- Find one or more entities, based on non-spatial aspects, then use spatial operations to get interesting data associated with these items
 - ▶ Eg find the length of the river called 'Mississippi'
 - Find the river based on name (non-spatial), use spatial operation to compute the length (assuming it is not stored as a numeric value)
 - ▶ Eg find the total area of all counties whose population exceeds 1,000,000
 - Find the counties with population exceeds 1,000,000 (non-spatial), compute each county's area(spatial operation) and sum all up.
- To answer these: use conventional index or table scan to find the appropriate rows, then call spatial functions on the spatial attribute of each



Spatial selection queries

- Find items whose spatial attribute has certain properties
 - ▶ Eg find rivers whose length is at least 10000
- **WHERE** clause will involve spatial operations
- Typical processing: scan all rows, apply appropriate spatial operation to the spatial attribute of each

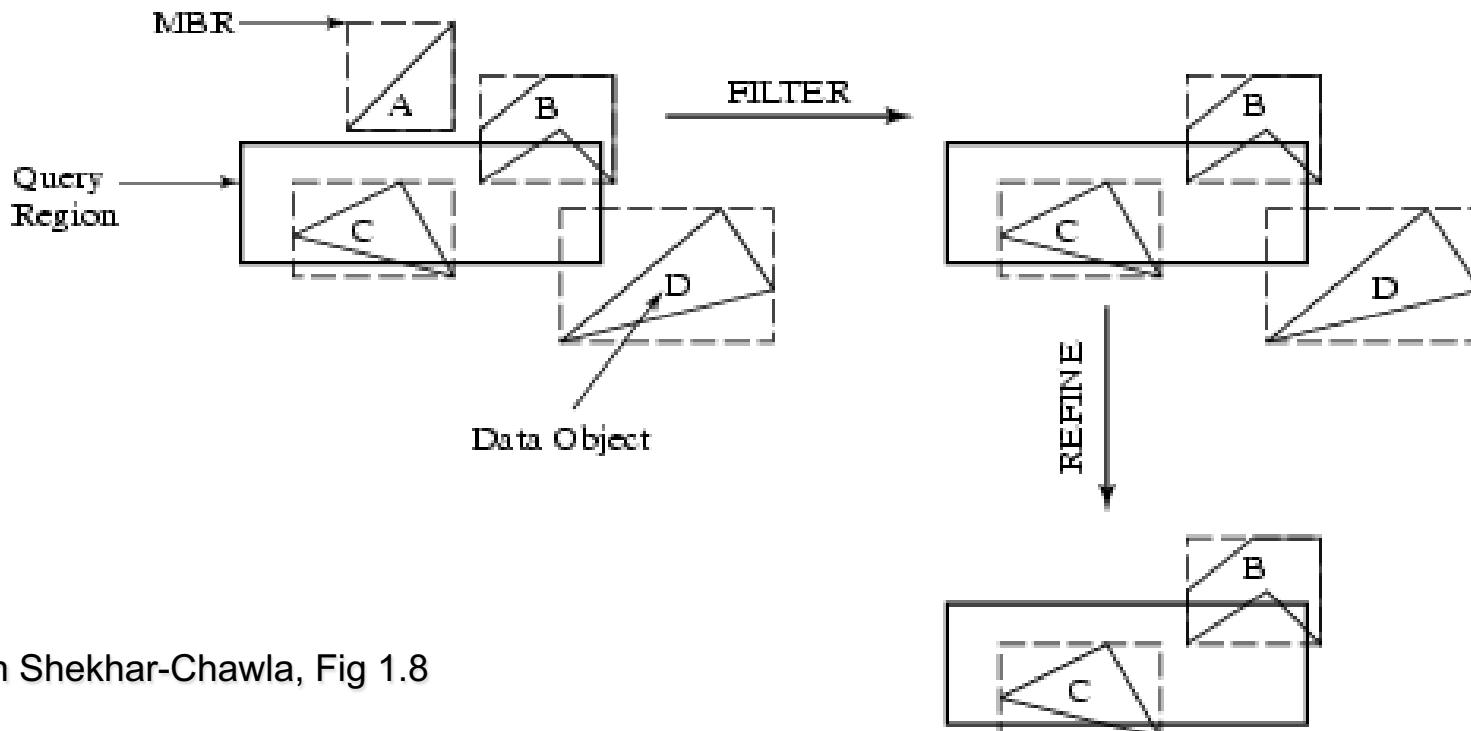


Spatial Range Query

- A particular kind of spatial selection, in which the condition involves a topological or metric relationship to a given object
 - ▶ Eg find all bookshops whose location is inside a given region
 - ▶ Eg find all farms that contain (part of) a given curve
 - ▶ Eg find all rivers that flow through a given region
 - ▶ Eg find all bookshops within 100 km of a given point
 - Equivalent to: find all whose location is inside a circular region of radius 100 km!
- Simple processing: scan the appropriate table, apply appropriate spatial operation to each item's spatial attribute
- But often one can do better: first filter to find a small set where the condition might be feasible; then refine the list by checking in detail each that pass the filter

Spatial Query Processing: Filter-Refine Strategy

- Eg find objects that intersect a query region
- Filter Step:** made easy if each object has associated to it a simple shape that surrounds it (eg Minimum Bounding Rectangle)
 - If object's MBR doesn't intersect query {or MBR of query}, there is no possibility that the object itself will intersect the query
- Refine Step:** Actually perform intersection method for those objects that get through the filter



From Shekhar-Chawla, Fig 1.8



Nearest neighbours

- Find entities that are as close as possible to given location,
Always give a bound on how many to find (K Nearest Neighbours)
 - ▶ Eg find 5 closest restaurants to (100, 350) and return them ranked by closeness
- Simple processing: scan all, compute distance; keep track of the ones with lowest distances seen so far
- Many index based algorithms, see next week



Spatial Join Query

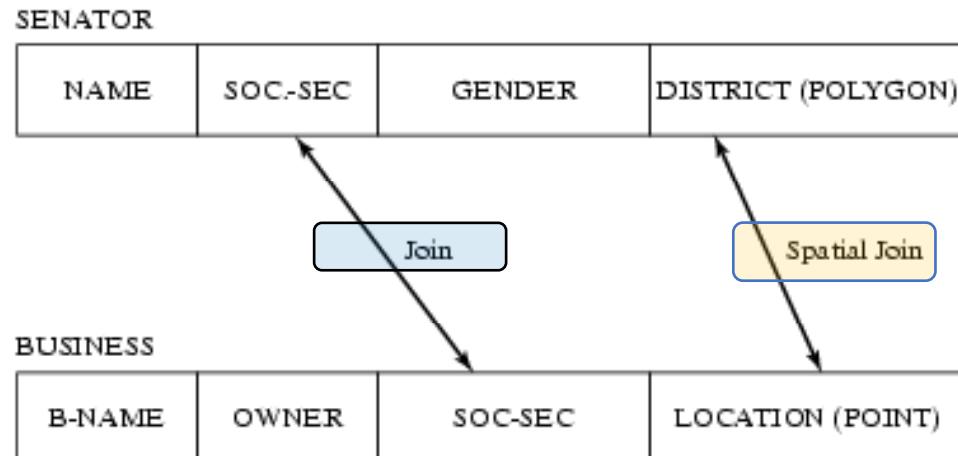
Spatial join example

```
SELECT      S.name  
FROM        Senator S, Business B  
WHERE       S.district.Area() > 300  
           AND Within(B.location, S.district)
```

Non-Spatial Join example

```
SELECT S.name  
FROM   Senator S, Business B  
WHERE  S.soc-sec = B.soc-sec AND S.gender = 'Female'
```

Similar to non-spatial join, spatial join are usually very expensive to process



From Shekhar-Chawla, Fig 1.7



Outline

- Motivation
- Spatial Data Model concepts
- Spatial Data Queries
- Spatial Query in MongoDB



Spatial data: GeoJSON

- Spatial data in MongoDB can be stored as **GeoJSON** object or as legacy coordinate pairs
 - ▶ **GeoJSON** data assumes earth-like sphere
 - ▶ Legacy coordinate pairs assumes flat plane
- **GeoJSON** object uses JSON format to represent spatial objects in OpenGIS
 - ▶ Point, LineString, Polygon, MultiPoint, MultiLineString, MultiPolygon, Geometry Collection
 - ▶ General format

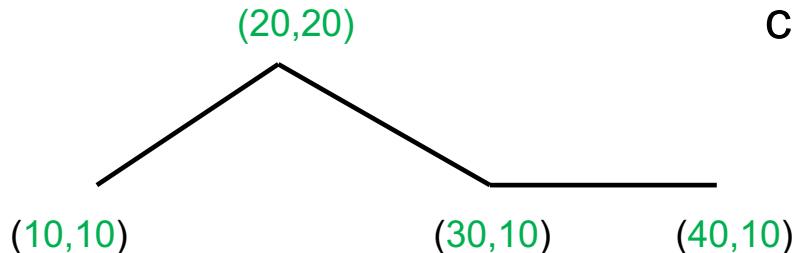
```
{ type: "<GeoJSON type>" , coordinates: <coordinates> }
```
 - ▶ The coordinate reference system for all GeoJSON coordinates is a geographic coordinate reference system, using the World Geodetic System 1984 (WGS 84) [WGS84] datum, with longitude and latitude units of decimal degrees.



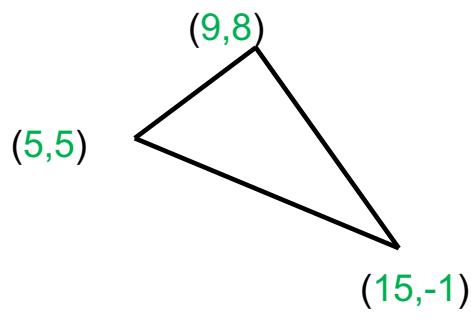
Point and LineString



{type: "Point", coordinates: [40,5]}



{type: "LineString",
coordinates: [
[10,10], [20,20], [30,10], [40,10]] }



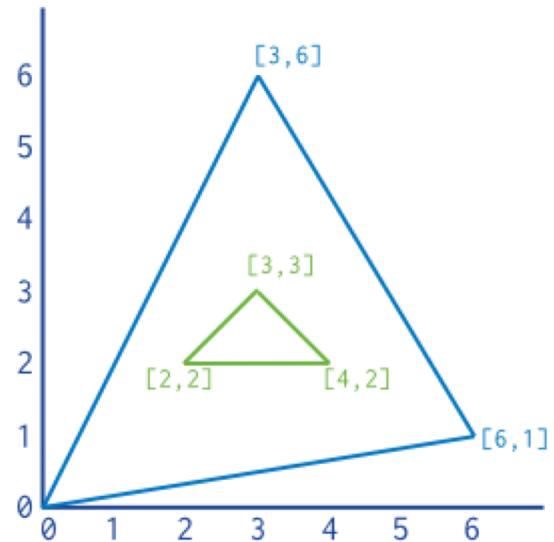
A closed LineString is called a linear ring, with at least four coordinate pairs and specify the same position as the first and last coordinates.

{type: "LineString",
coordinates: [[5,5],[9,8], [15,-1], [5,5]]
}

Polygon

- Polygon is used to model two dimensional surface
 - ▶ Triangle, Rectangle, Pentagon, ...
 - ▶ A polygon is represented as one or many linear rings, the first is the exterior ring bounds the surface, the others are interior rings bound holes within the surface

```
{  
  type : "Polygon",  
  coordinates : [  
    [[ 0 , 0 ],[ 3 , 6 ],[ 6 , 1 ],[ 0 , 0 ]],  
    [[ 2 , 2 ],[ 3 , 3 ],[ 4 , 2 ],[ 2 , 2 ]]  
  ]  
}
```



Model multiple disjoint objects

■ MultiPoints, MultiPolygon

```
{ type: "MultiPoint",
  coordinates: [ [ -73.9580, 40.8003 ],
    [ -73.9498, 40.7968 ],
    [ -73.9737, 40.7648 ],
    [ -73.9814, 40.7681 ] ] }

{ type: "MultiPolygon",
  coordinates: [ [ [ [ -73.95, 40.80 ], [ -73.9498, 40.79 ], [ -73.97, 40.76 ], [ -73.95, 40.80 ] ] ],
    [ [ [ -73.95, 40.80 ], [ -73.94, 40.79 ], [ -73.97, 40.76 ], [ -73.95, 40.80 ] ] ] ] }
```

■ Geometry collection

```
{ type: "GeometryCollection",
  geometries: [ { type: "MultiPoint",
      coordinates: [ [ -73.95, 40.80 ], [ -73.94, 40.79 ] ] },
    { type: "MultiLineString",
      coordinates: [ ... ] }
  ] }
```



Spatial Queries

■ **\$near** and **\$nearSphere**

- ▶ Specifies a point for which a **geospatial** query returns the documents from nearest to farthest
- ▶ Can be used to run queries like find all car parks/restaurants within certain distance

■ **\$geoWithin**

- ▶ Find all geo objects contained in a query shape

■ **\$geoIntersects**

- ▶ Find all geo objects intersects with a query shape. Here intersect includes relationships such as cover, equal, overlap, touch and so on.

■ Others



Spatial Index

- **2dsphere** indexes supports all MongoDB geospatial queries
- Eg.

```
db.places.insert(  
 {  
   loc : { type: "Point", coordinates: [ -73.97, 40.77 ] },  
   name: "Central Park",  
   category : "Parks"  
 }  
)
```

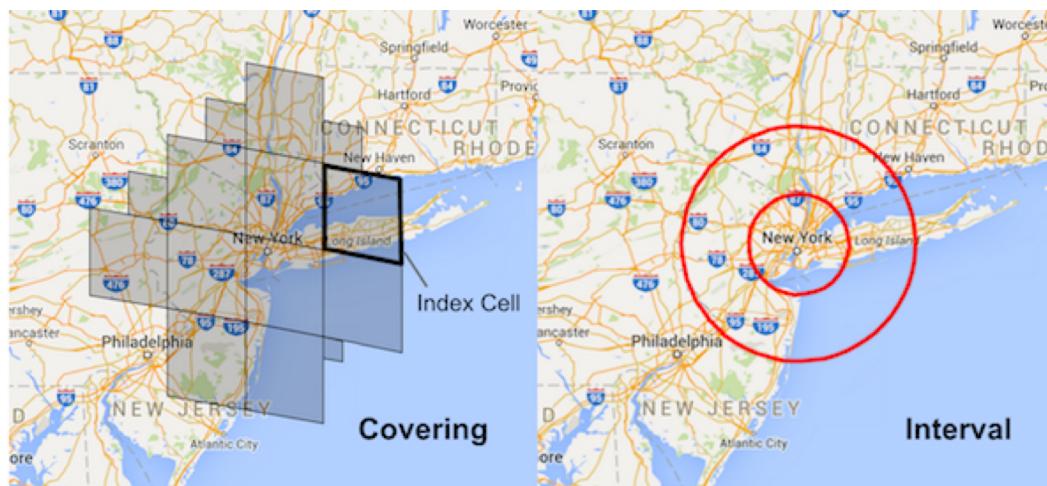
- **db.places.createIndex({ loc : "2dsphere" })**



MongoDB Spatial Index

- “ MongoDB’s 2dsphere index actually combines the strength of discrete global grids and B+ -tree structures, which first partitions the Earth surface into cells at multiple resolution levels and then applies a B+ -tree to index geographical features approximated as one or multiple cells.”

L. Xiang, J. Huang , X. Shao and D. Wang: MongoDB-Based Management of Planar Spatial Data with a Flattened R-Tree
<https://pdfs.semanticscholar.org/860f/0cfe3e4b4cb66b2b2895016a60e824e9e9f8.pdf>

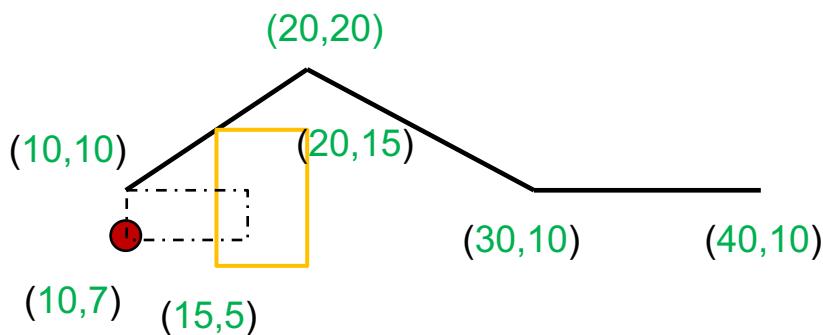


MongoDB Blog: Geospatial Performance Improvements in MongoDB 3.2
<https://www.mongodb.com/blog/post/geospatial-performance-improvements-in-mongodb-3-2>

Spatial Query -- \$geoIntersects

```
db.places.find({  
    loc :{$geoIntersects:  
        {$geometry:{  
            type: "Polygon",  
            coordinates: [[[10,7], [10,10], [17,10],[17,7],[10,7]]]}  
        }  
    }  
})
```

● (20,30)



References

- S. Shekhar and S.Chawla: *Spatial Databases: A Tour*. Prentice Hall, 2002. [<http://www.spatial.cs.umn.edu/Book/>]
 - ▶ Chapter 1-3
- MongoDB document on geospatial query
 - ▶ <https://docs.mongodb.com/manual/geospatial-queries/>

