



Motivation

- Many entities represent physical objects, and some physical features also matter for business purposes
 - ▶ Eg. A store has a name, business category, contact number and is located at a particular place (geo-spatial)
 - The geo-spatial feature can help find a store that is close to a customer
 - ▶ Eg. A toy has name, category, material, and shape
 - The shape feature helps to find a box that can fit the toy
- There are large amount of geospatial data:
 - ▶ Businesses and homes have addresses
 - Both the logic aspect and physical aspect
 - ▶ Google Maps, Google Earth
 - ▶ Weather and Climate Data

Outline

- Motivation
- Spatial Data Model
- Spatial Data Queries
- Spatial Query in MongoDB

Spatial Data and Object Concept

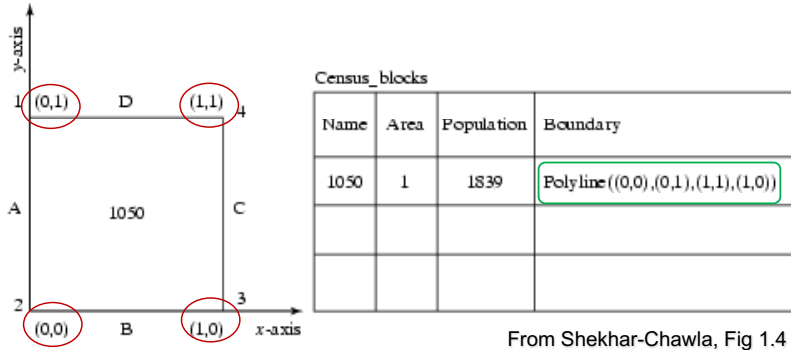
- Spatial feature could refer to a
 - ▶ Point, a 2d shape, a 3d shape, or shape in higher dimension
- Spatial features of an entity cannot be represented as simple value type
 - ▶ A point in 2d space has two coordinate values
- It is natural to use object to represent spatial features (spatial data)
- There are also spatial related operations we need to perform on spatial data
 - ▶ Compute distance between points
 - ▶ Compute area of 2d shape or volume of 3d shapes
 - ▶ Compute various relationships among spatial objects

Spatial Data as Object Example

- Consider a spatial data representing census block:

```
CREATE TABLE census_blocks (
  name      string,
  area      float,
  population number,
  boundary  Polyline );
```

A User
Defined Type



Spatial Data in Purely Relational Form

Census_blocks				Polygon	
Name	Area	Population	boundary-ID	boundary-ID	edge-name
340	1	1839	1050	1050	A
				1050	B
				1050	C
				1050	D

Edge	
edge-name	endpoint
A	1
A	2
B	2
B	3
C	3
C	4
D	4
D	1

Point		
endpoint	x-coor	y-coor
1	0	1
2	0	0
3	1	0
4	1	1

From Shekhar-Chawla, Fig 1.4

Spatial Database Management System

- A SDBMS is a software module that
 - can work with an underlying DBMS
 - supports spatial data models, spatial abstract data types (ADTs) and a query language from which these ADTs are callable
 - supports spatial indexing, efficient algorithms for processing spatial operations, and domain specific rules for query optimization
 - Many RDBMS and NoSQL storage systems have support for spatial data
 - Oracle, SQL Server, MongoDB, Neo4j, ...
- SDBMS components
 - spatial data model
 - query language
 - query processing
 - file organization and indices
 - query optimization
 - etc.

Outline

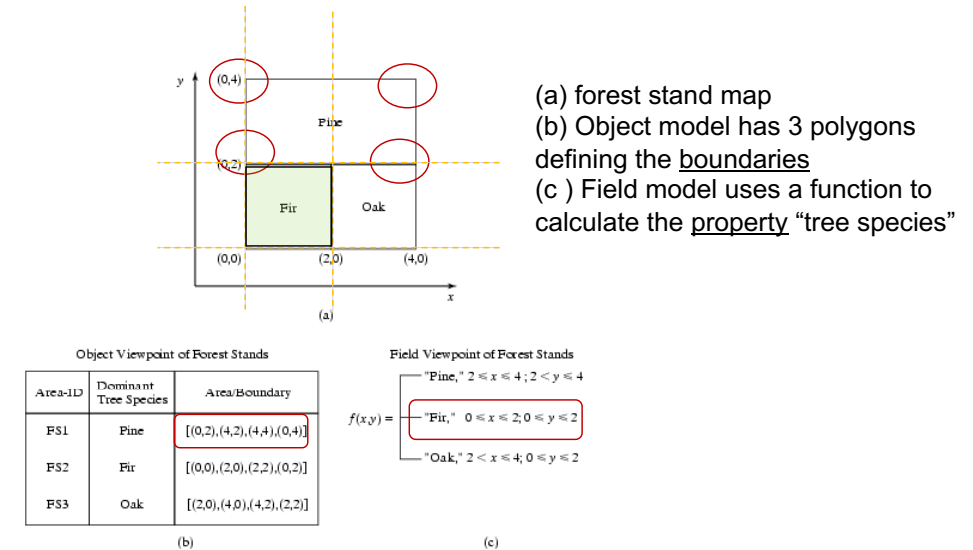
- Motivation
- Spatial Data Model
 - Field vs. Object Models
 - Coordinate System
 - Topological Operations
- Spatial Data Queries
- Spatial Query in MongoDB

Models of Spatial Information

Two common models

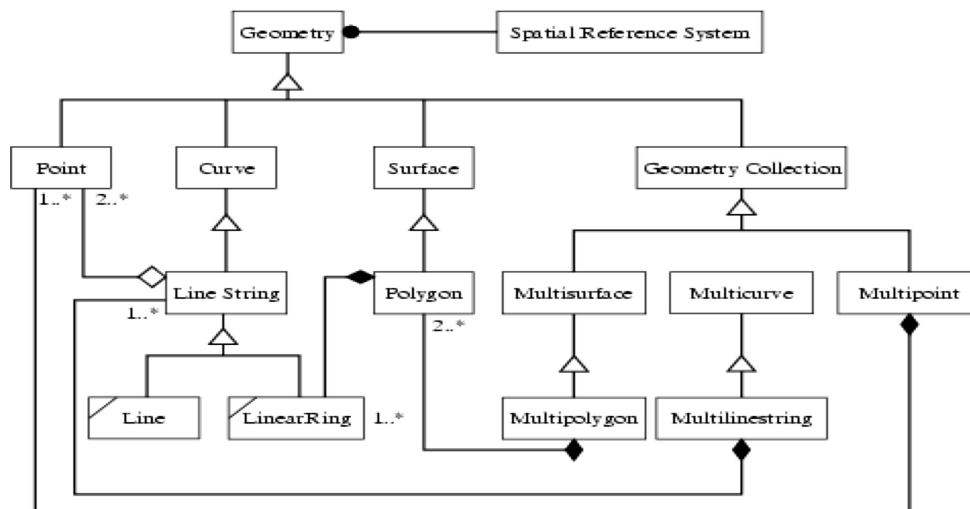
- Field based (*also: space-based*)
 - Model properties of underlying space
 - Good for expressing values vary continuously over space (e.g. temperature, rainfall, elevation, depth, etc.
 - Fields are actually **functions** that map spatial locations to values
- Object based (*also: entity-based*)
 - Model boundaries of spatial feature (e.g. the census block is modelled by a polygon)
 - Good for expressing discrete spatial entities

Examples of Field and Object Models



From Shekhar-Chawla, Fig 2.1

OpenGIS Geometry Model




From Shekhar-Chawla, Fig 2.2

Example Spatial Objects

- **Point:** represented by its coordinates eg (-10, 30)
- Collection of several **points**
- **Line String**
 - Simplest form is piece-wise linear, given by points (and implying the straight segments between them)
 - Eg (0,1), (1,1), (2,2)
- **Polygon**
 - A 2-d region whose boundary is given
 - Simplest form: boundary is a **line string** that returns to its start
 - More complicated: region with holes
- Collection of **polygons**

Coordinate Systems

- Points from a 2-d space are represented by pairs of numbers
 - ▶ The numbers could refer to a dot on a drawing area, a piece of land in a game setting, or a location on earth
- There are many ways to associate numbers with points
- Simple 2d Cartesian coordinate
 - ▶ Choose a point as origin (0,0)
 - ▶ Choose a direction for the x-axis, and a scale (how far is (1,0)) from (0,0)?)



Operations on Spatial Objects in the Object Model

- Classifying operations
 - ▶ **Set based:**
 - a set operation (e.g. intersection) of 2 polygons produce another polygon
 - ▶ **Topological operations:** Boundary of USA touches boundary of Canada
 - ▶ **Directional:** New York city is to east of Chicago
 - ▶ **Metric:** Chicago is about 700 miles from New York city.

Set theory based	Union, Intersection, Containment
Topological	Touches, Disjoint, Overlap, etc.
Directional	East, North-West, etc.
Metric	Distance

A Round World

- The surface of the earth is 2-dimensional, but curved
 - ▶ Cartesian systems work reasonably in small regions
- Traditional geographic coordinate system
 - ▶ 2d: longitude and latitude
 - ▶ 3d: longitude, latitude, elevation
 - ▶ The surface is a sphere
 - (-179,10) is very close to (179, 10)
 - A linestring might cross the (long = 180) line;
- Many SDBMS supports both flat space and sphere


Topological Relationships

- Topological Relationships
 - ▶ invariant under elastic deformation (without tear, merge).
 - ▶ Two countries which touch each other in a planar paper map will continue to do so in spherical globe maps.
- Example queries with topological operations
 - ▶ What is the topological relationship between two objects A and B ?
 - ▶ Find all objects which have a given topological relationship to object A?
 - E.g. find all rivers that cross a city
- Can we express *topological relationship* mathematically?
 - ▶ Metric operations may be expressed using various functions, e.g. distance function
 - ▶ Set operations can be express mathematically
 - ▶ The mathematical form helps to calculate such relationships

Topological Concepts

■ Interior, boundary, exterior

- ▶ Let A be an object in a "Universe" U .



Green is A interior (A^o)

Red is boundary of A (∂A)

Blue – (Green + Red) is A exterior (A^-)

- ▶ Exterior is also referred to as the *complement* of an object

Nine-Intersection Model of Topological Relationships

■ Many topological Relationship between A and B can be specified using 9 intersection model

- ▶ Eight possible 2D topological relationships for objects without holes;

■ Nine intersections

- ▶ intersections between interior, boundary, exterior of A , B
- ▶ A and B are spatial objects in a two dimensional plane.
- ▶ Can be arranged as a 3 by 3 matrix
- ▶ Matrix element take a value of 0 (false) or 1 (true)

$$\Gamma_9(A, B) = \begin{pmatrix} A^0 \cap B^0 & A^0 \cap \partial B & A^0 \cap B^- \\ \partial A \cap B^0 & \partial A \cap \partial B & \partial A \cap B^- \\ A^- \cap B^0 & A^- \cap \partial B & A^- \cap B^- \end{pmatrix}$$

From Shekhar-Chawla, p 28

Specifying Topological Operations using the 9-Intersection Model

		B		
		I	B	E
A	I			
	B			
	E			

$\begin{pmatrix} 0 & 0 & 1 \\ 0 & 0 & 1 \\ 1 & 1 & 1 \end{pmatrix}$	$\begin{pmatrix} 1 & 1 & 1 \\ 0 & 0 & 1 \\ 0 & 0 & 1 \end{pmatrix}$	$\begin{pmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 1 & 1 & 1 \end{pmatrix}$	$\begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$
disjoint	contains	inside	equal
$\begin{pmatrix} 0 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}$	$\begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{pmatrix}$	$\begin{pmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 1 & 1 & 1 \end{pmatrix}$	$\begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}$
meet	covers	coveredBy	overlap

For **disjoint** relation:
A's exterior intersects with B's everything and vice versa

If A **contains** B:
A's interior intersects with B's everything and B's exterior intersects with A's everything

From Shekhar-Chawla, Fig 2.3

Outline

■ Motivation

■ Spatial Data Model concepts

■ Spatial Data Queries

- ▶ Query type
- ▶ General processing steps

■ Spatial Query in MongoDB

Spatial Processing

- Find one or more entities, based on non-spatial aspects, then use spatial operations to get interesting data associated with these items
 - Eg find the length of the river called 'Mississippi'
 - Find the river based on name (non-spatial), use spatial operation to compute the length (assuming it is not stored as a numeric value)
 - Eg find the total area of all counties whose population exceeds 1,000,000
 - Find the counties with population exceeds 1,000,000 (non-spatial), compute each county's area(spatial operation) and sum all up.
- To answer these: use conventional index or table scan to find the appropriate rows, then call spatial functions on the spatial attribute of each

Spatial selection queries


- Find items whose spatial attribute has certain properties
 - Eg find rivers whose length is at least 10000
- WHERE** clause will involve spatial operations
- Typical processing: scan all rows, apply appropriate spatial operation to the spatial attribute of each

Spatial Range Query

- A particular kind of spatial selection, in which the condition involves a topological or metric relationship to a given object
 - Eg find all bookshops whose location is inside a given region
 - Eg find all farms that contain (part of) a given curve
 - Eg find all rivers that flow through a given region
 - Eg find all bookshops within 100 km of a given point
 - Equivalent to: find all whose location is inside a circular region of radius 100 km!
- Simple processing: scan the appropriate table, apply appropriate spatial operation to each item's spatial attribute
- But often one can do better: first filter to find a small set where the condition might be feasible; then refine the list by checking in detail each that pass the filter

Spatial Query Processing: Filter-Refine Strategy

- Eg find objects that intersect a query region
- Filter Step:** made easy if each object has associated to it a simple shape that surrounds it (eg Minimum Bounding Rectangle)
 - If object's MBR doesn't intersect query {or MBR of query}, there is no possibility that the object itself will intersect the query
- Refine Step:** Actually perform intersection method for those objects that get through the filter



From Shekhar-Chawla, Fig 1.8

Nearest neighbours

- Find entities that are as close as possible to given location, Always give a bound on how many to find (K Nearest Neighbours)
 - Eg find 5 closest restaurants to (100, 350) and return them ranked by closeness
- Simple processing: scan all, compute distance; keep track of the ones with lowest distances seen so far
- Many index based algorithms, see next week

Spatial Join Query

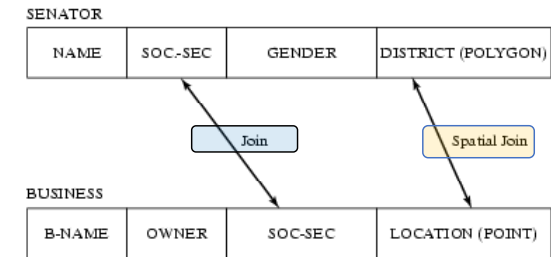
Spatial join example

```
SELECT S.name
FROM Senator S, Business B
WHERE S.district.Area() > 300
AND Within(B.location, S.district)
```

Non-Spatial Join example

```
SELECT S.name
FROM Senator S, Business B
WHERE S.soc-sec = B.soc-sec AND S.gender = 'Female'
```

- Similar to non-spatial join, spatial join are usually very expensive to process



From Shekhar-Chawla, Fig 1.7

Outline


- Motivation
- Spatial Data Model concepts
- Spatial Data Queries
- Spatial Query in MongoDB

Spatial data: GeoJSON

- Spatial data in MongoDB can be stored as **GeoJSON** object or as legacy coordinate pairs
 - GeoJSON data assumes *earth-like* sphere
 - Legacy coordinate pairs assumes flat plane
- GeoJSON object uses JSON format to represent spatial objects in OpenGIS
 - Point, LineString, Polygon, MultiPoint, MultiLineString, MultiPolygon, Geometry Collection
 - General format

```
{ type: "<GeoJSON type>" , coordinates: <coordinates> }
```
 - The coordinate reference system for all GeoJSON coordinates is a geographic coordinate reference system, using the World Geodetic System 1984 (WGS 84) [WGS84] datum, with longitude and latitude units of decimal degrees.


Point and LineString



Polygon

- Polygon is used to model two dimensional surface
 - ▶ Triangle, Rectangle, Pentagon, ...
 - ▶ A polygon is represented as one or many linear rings, the first is the exterior ring bounds the surface, the others are interior rings bound holes within the surface

```
{
  type: "Polygon",
  coordinates: [
    [[ 0, 0 ], [ 3, 6 ], [ 6, 1 ], [ 0, 0 ]],
    [[ 2, 2 ], [ 3, 3 ], [ 4, 2 ], [ 2, 2 ] ]
  ]
}
```



Model multiple disjoint objects

■ MultiPoints, MultiPolygon

```
{ type: "MultiPoint",
  coordinates: [ [ -73.9580, 40.8003 ],
                 [ -73.9498, 40.7968 ],
                 [ -73.9737, 40.7648 ],
                 [ -73.9814, 40.7681 ] ] }

{ type: "MultiPolygon",
  coordinates: [ [ [ [ -73.95, 40.80 ], [ -73.9498, 40.79 ], [ -73.97, 40.76 ], [ -73.95, 40.80 ] ] ],
                 [ [ [ -73.95, 40.80 ], [ -73.94, 40.79 ], [ -73.97, 40.76 ], [ -73.95, 40.80 ] ] ] ] }
```

■ Geometry collection

```
{ type: "GeometryCollection",
  geometries: [ { type: "MultiPoint",
                  coordinates: [ [ -73.95, 40.80 ], [ -73.94, 40.79 ] ] },
                { type: "MultiLineString",
                  coordinates: [ ... ] }
                ] }
```

Spatial Queries

■ \$near and \$nearSphere

- ▶ Specifies a point for which a **geospatial** query returns the documents from nearest to farthest
- ▶ Can be used to run queries like find all car parks/restaurants within certain distance

■ \$geoWithin

- ▶ Find all geo objects contained in a query shape

■ \$geoIntersects

- ▶ Find all geo objects intersects with a query shape. Here intersect includes relationships such as cover, equal, overlap, touch and so on.

■ Others

Spatial Index

- **2dsphere** indexes supports all MongoDB geospatial queries

- **Eg.**


```
db.places.insert(
{
  loc : { type: "Point", coordinates: [ -73.97, 40.77 ] },
  name: "Central Park",
  category : "Parks"
}
```

- `db.places.createIndex({ loc : "2dsphere" })`

MongoDB Spatial Index

- “ MongoDB’s 2dsphere index actually combines the strength of discrete global grids and B+ -tree structures, which first partitions the Earth surface into cells at multiple resolution levels and then applies a B+ -tree to index geographical features approximated as one or multiple cells.”


L. Xiang, J. Huang, X. Shao and D. Wang: MongoDB-Based Management of Planar Spatial Data with a Flattened R-Tree
<https://pdfs.semanticscholar.org/860f/0cfe3e4b4cb66b2b2895016a60e824e9e9f8.pdf>



MongoDB Blog: Geospatial Performance Improvements in MongoDB 3.2
<https://www.mongodb.com/blog/post/geospatial-performance-improvements-in-mongodb-3-2>

Spatial Query -- \$geoIntersects

```
db.places.find({
  loc :{$geoIntersects:
    {$geometry:{
      type: "Polygon",
      coordinates:[[[10,7], [10,10], [17,10],[17,7],[10,7]]]}
    }
  })
```



References

- S. Shekhar and S.Chawla: *Spatial Databases: A Tour*. Prentice Hall, 2002. [<http://www.spatial.cs.umn.edu/Book/>]
 - ▶ Chapter 1-3
- MongoDB document on geospatial query
 - ▶ <https://docs.mongodb.com/manual/geospatial-queries/>